

# Exploring The Power Of Conv'Nets

## 01 Introduction to Computer Vision

[Yahia Chammami](#)



# What is Machine Learning?

Machine learning is a subdomain of computer science that focuses on algorithms which help a computer learn from data without explicit programming



# Types of Machine Learning

01

**Supervised learning** - uses labeled inputs (meaning the input has a corresponding output label) to train models and learn outputs

02

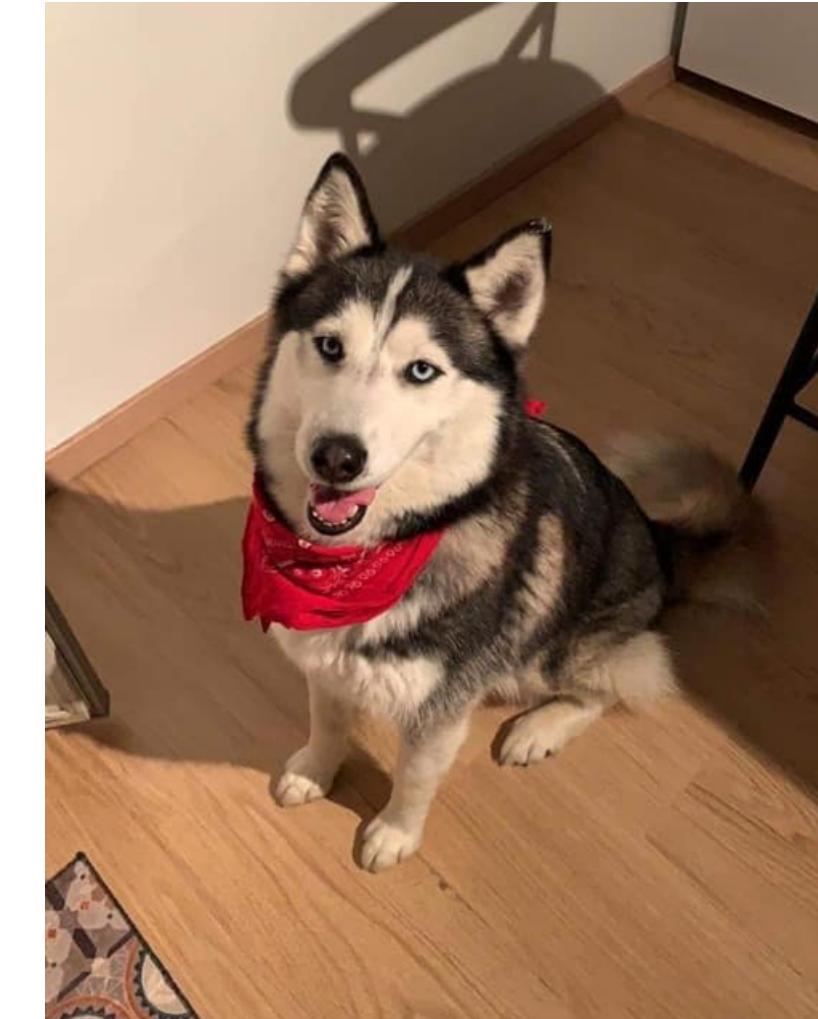
**Unsupervised learning** - uses unlabeled data to learn about patterns in data

03

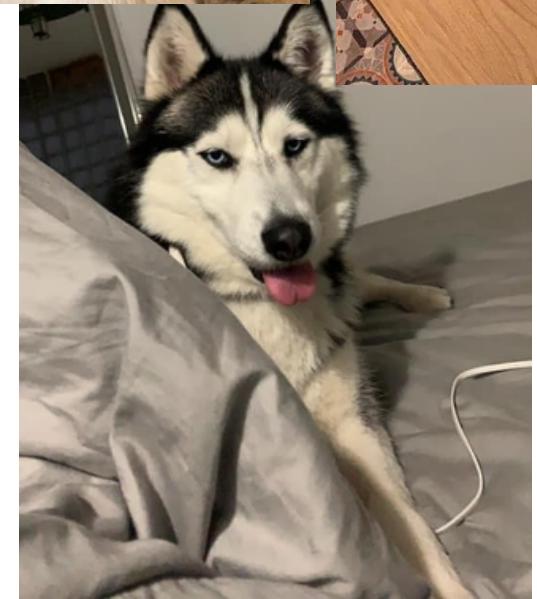
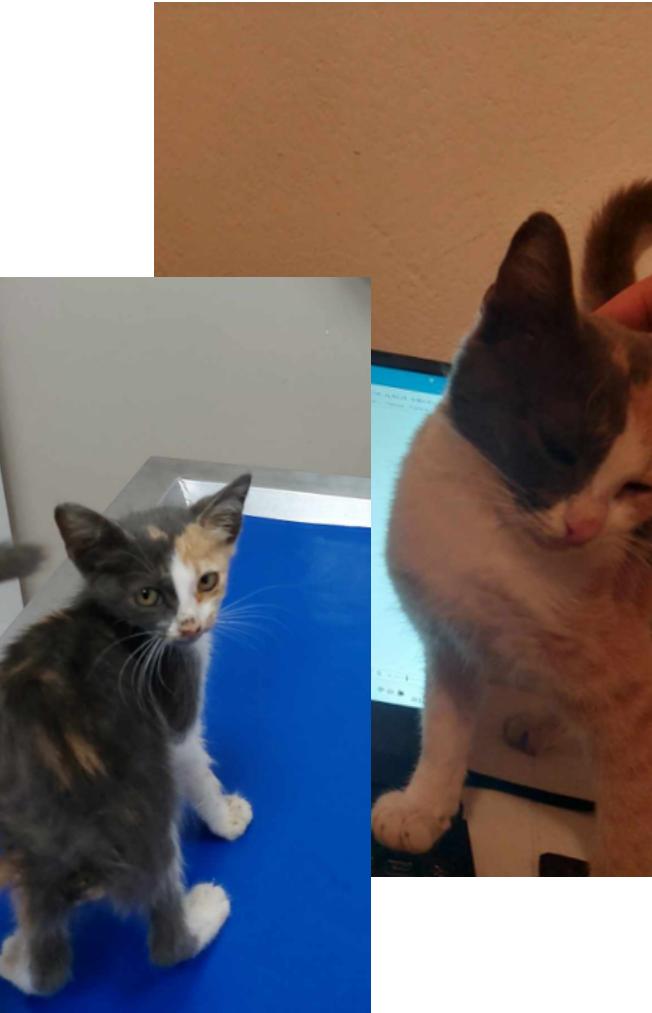
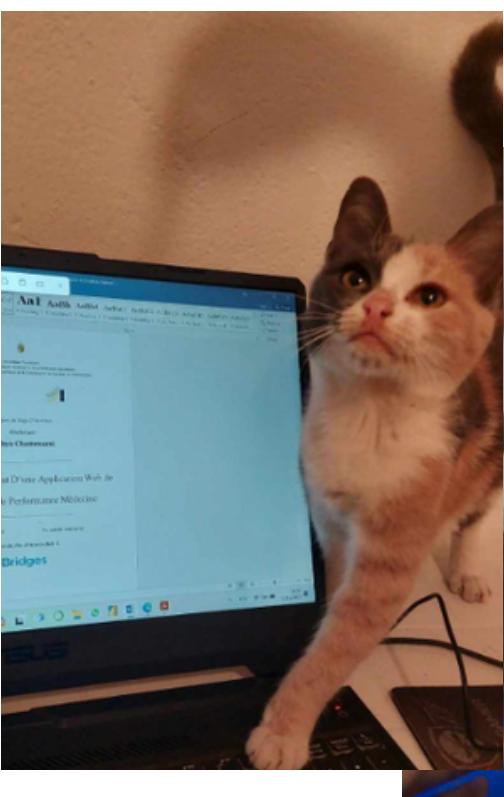
**Reinforcement learning** - agent learning in interactive environment based on rewards and penalties



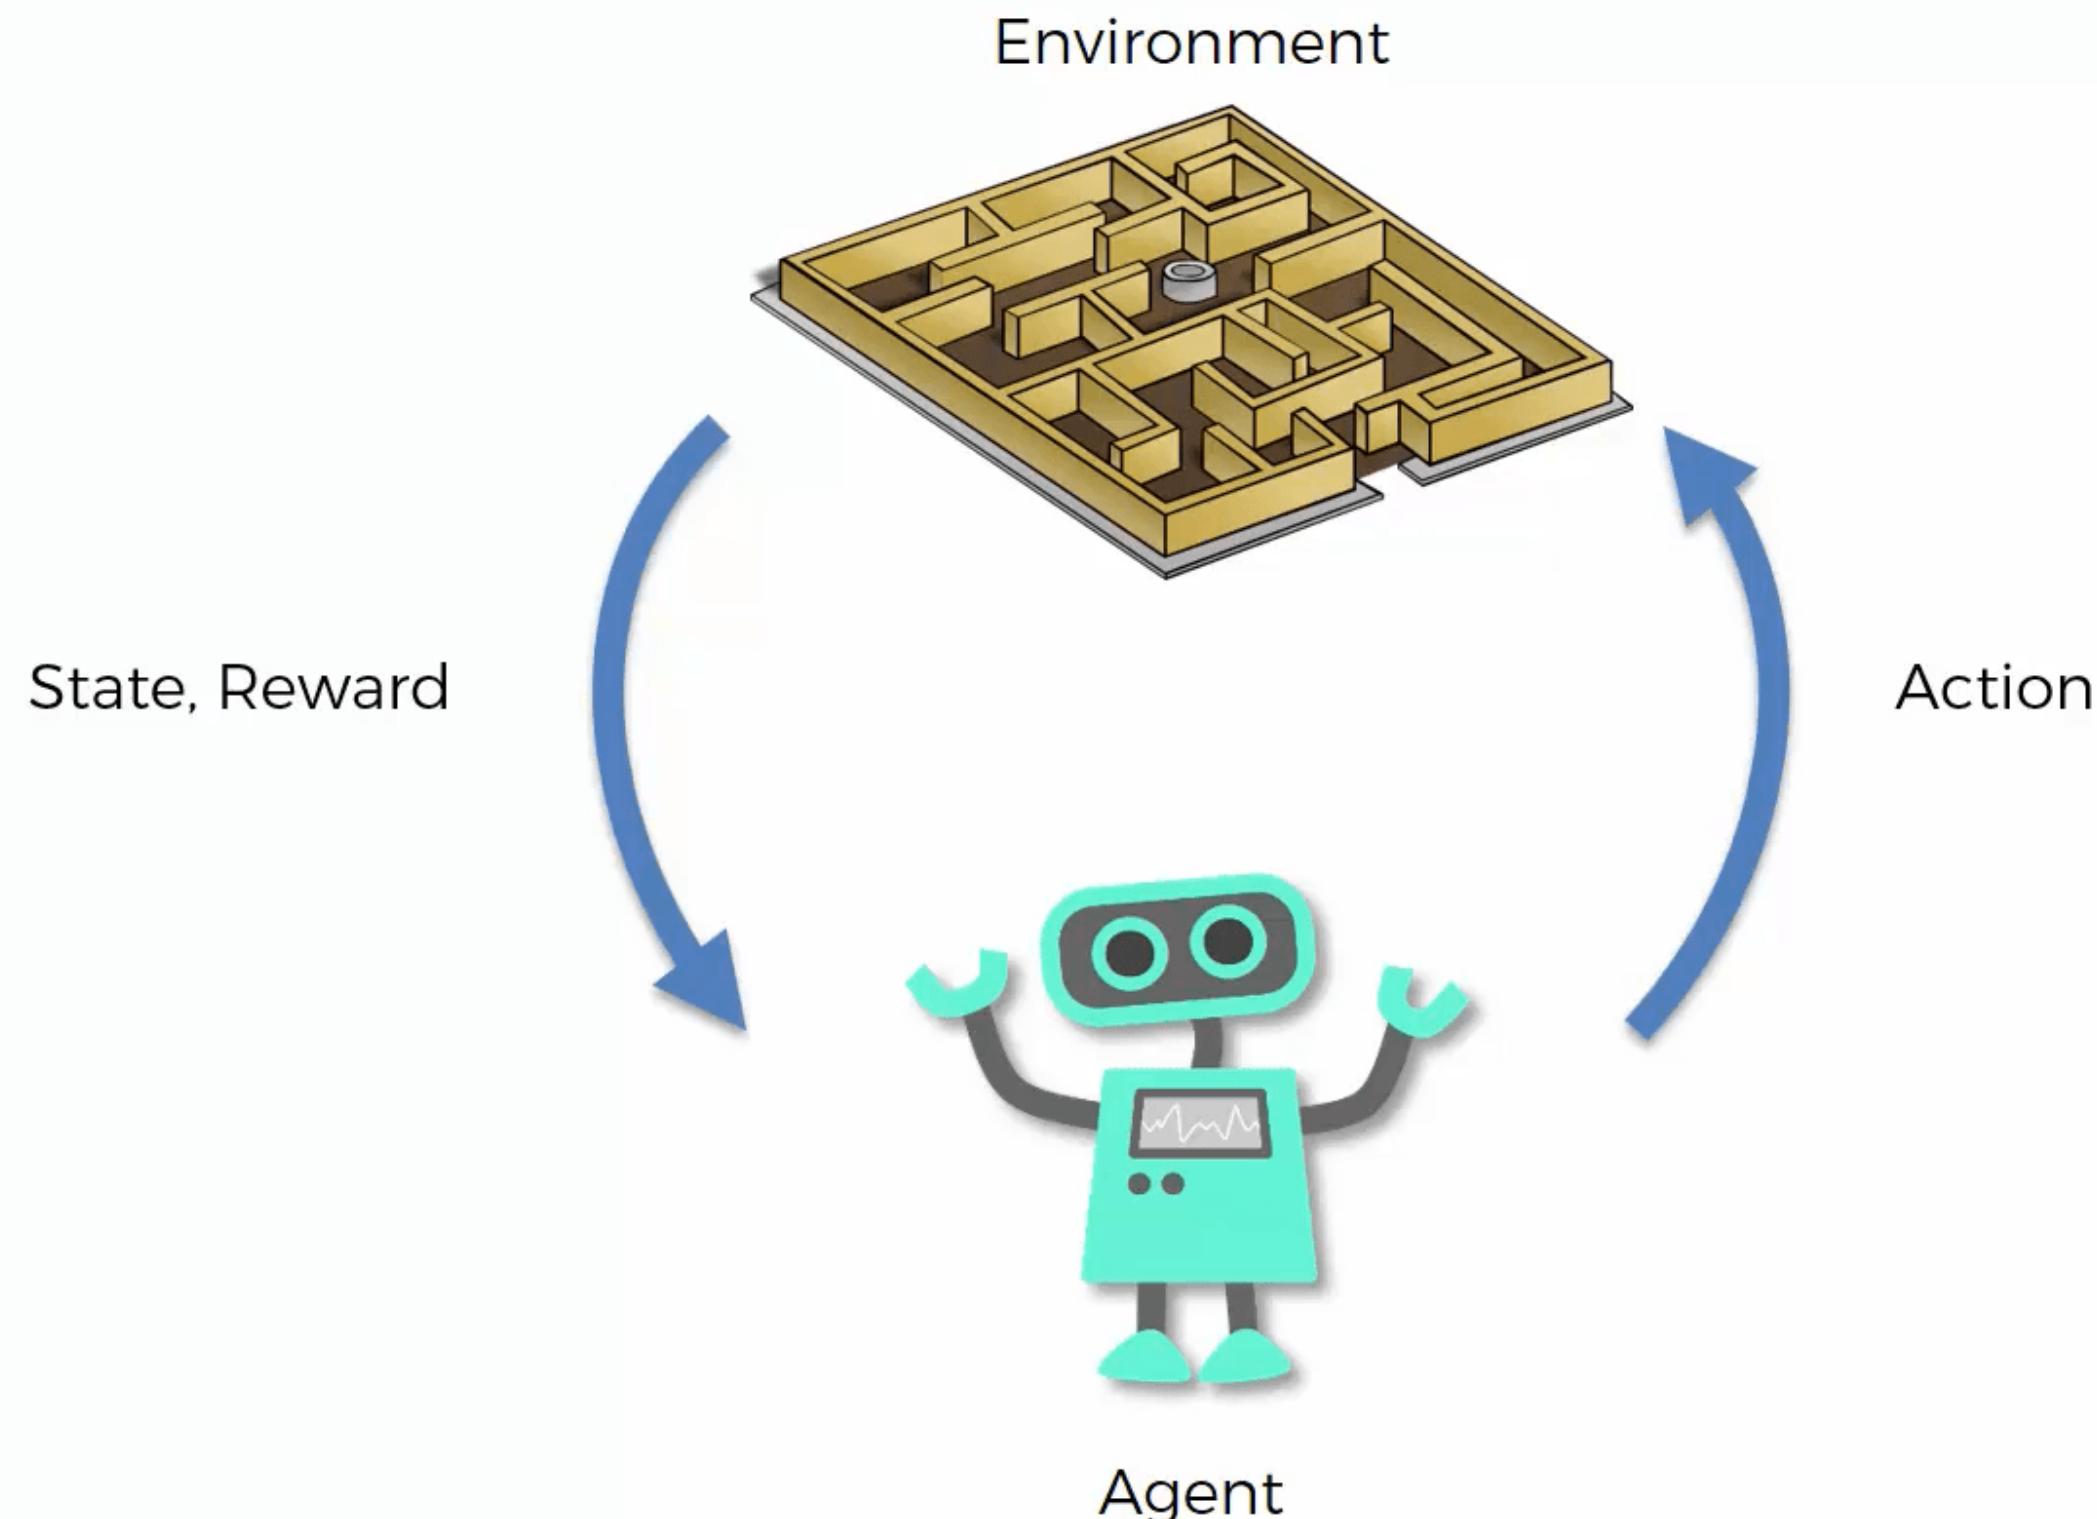
# Supervised Learning



# UnSupervised Learning



# Reinforcement learning



# Supervised Learning Tasks

01 Regression - predict continuous values

02 Classification - predict discrete classes



## BINARY CLASSIFICATION

Positive/negative

Cat/dog

Spam/not spam



## MULTICLASS CLASSIFICATION

Cat/dog/lizard/dolphin

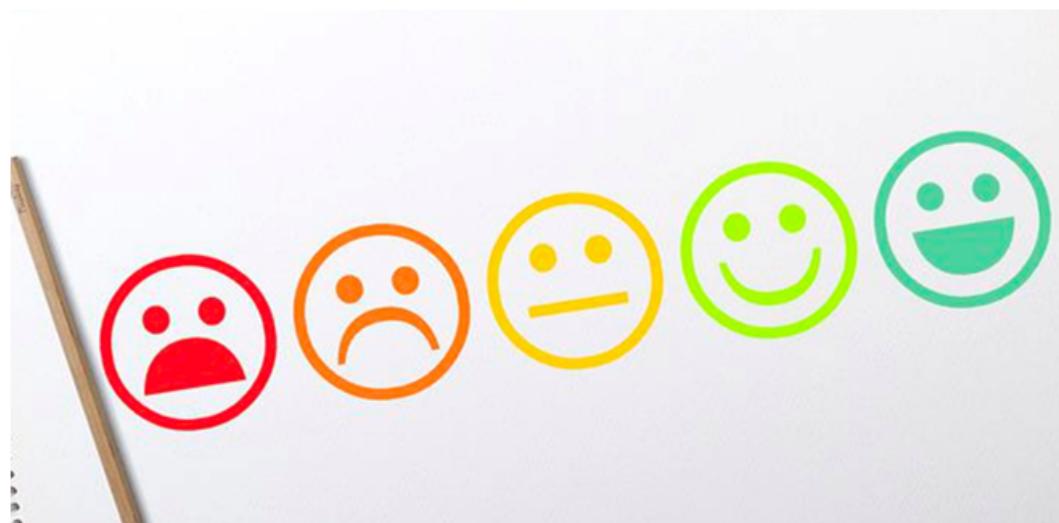
Orange/apple/pear

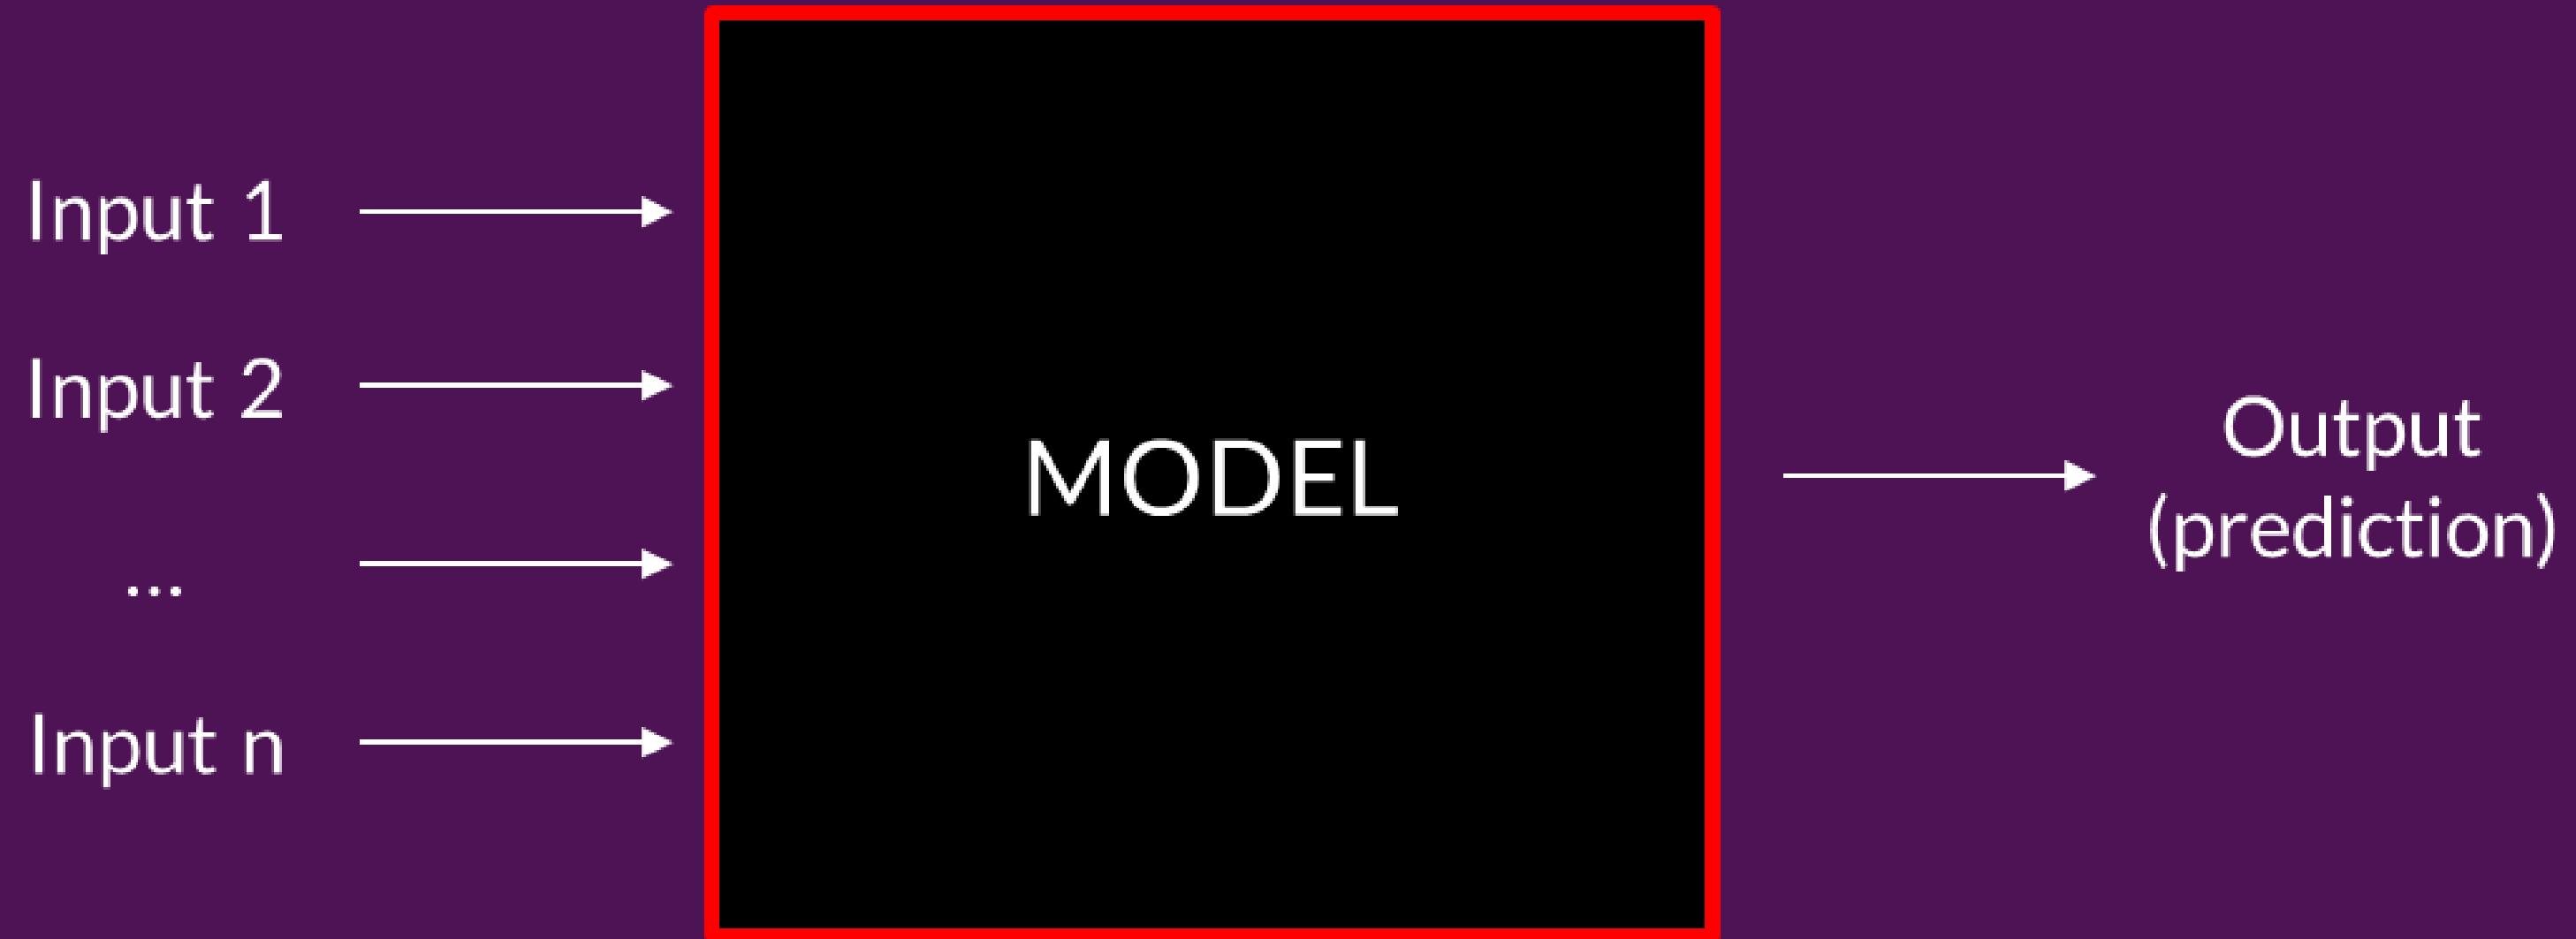
Plant species

# **Quantitative - numerical valued data (could be discrete or continuous)**



# **Qualitative - categorical data (finite number of categories or groups)**

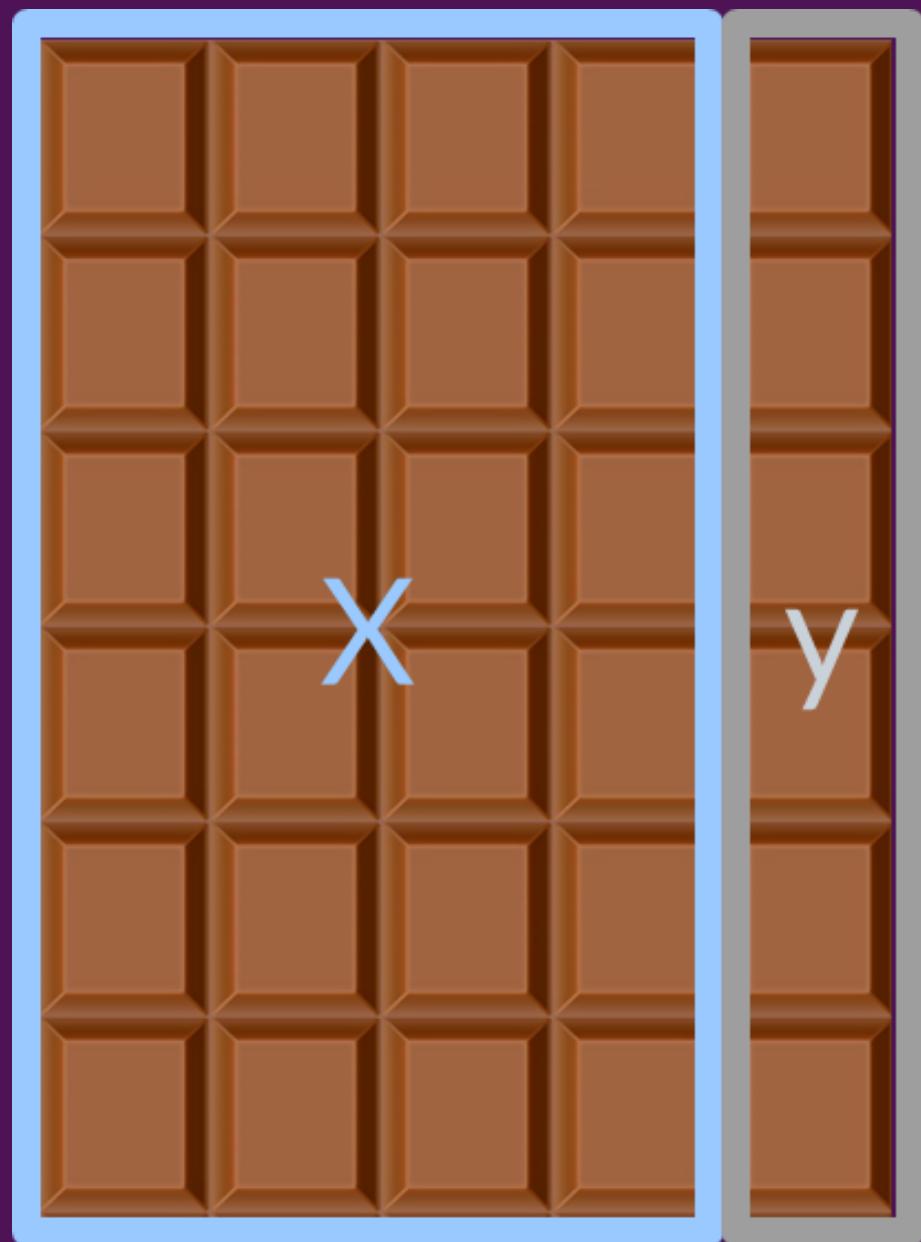




How do we make the model learn?  
How can we tell whether or not it's learning?



# Supervised Learning Dataset



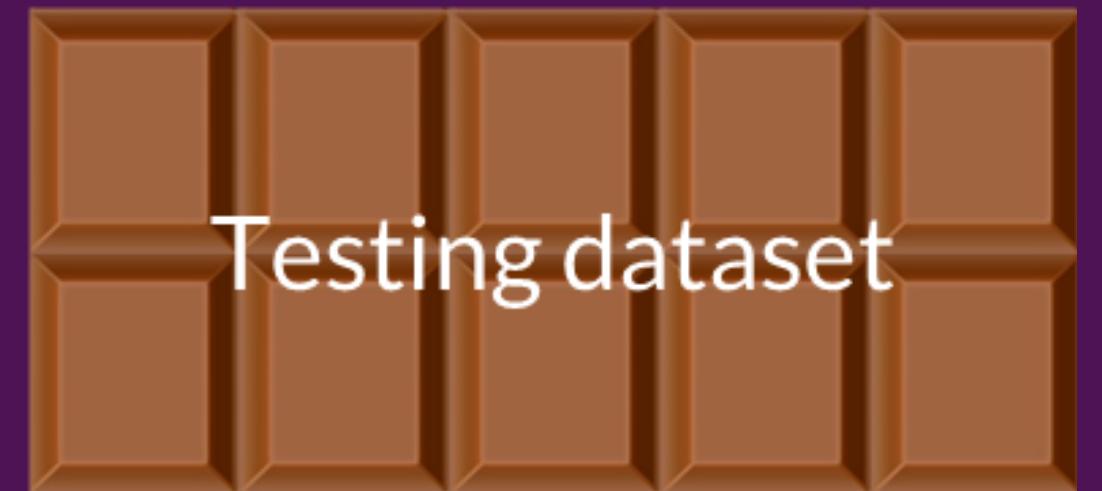
# Supervised Learning Dataset



Training dataset



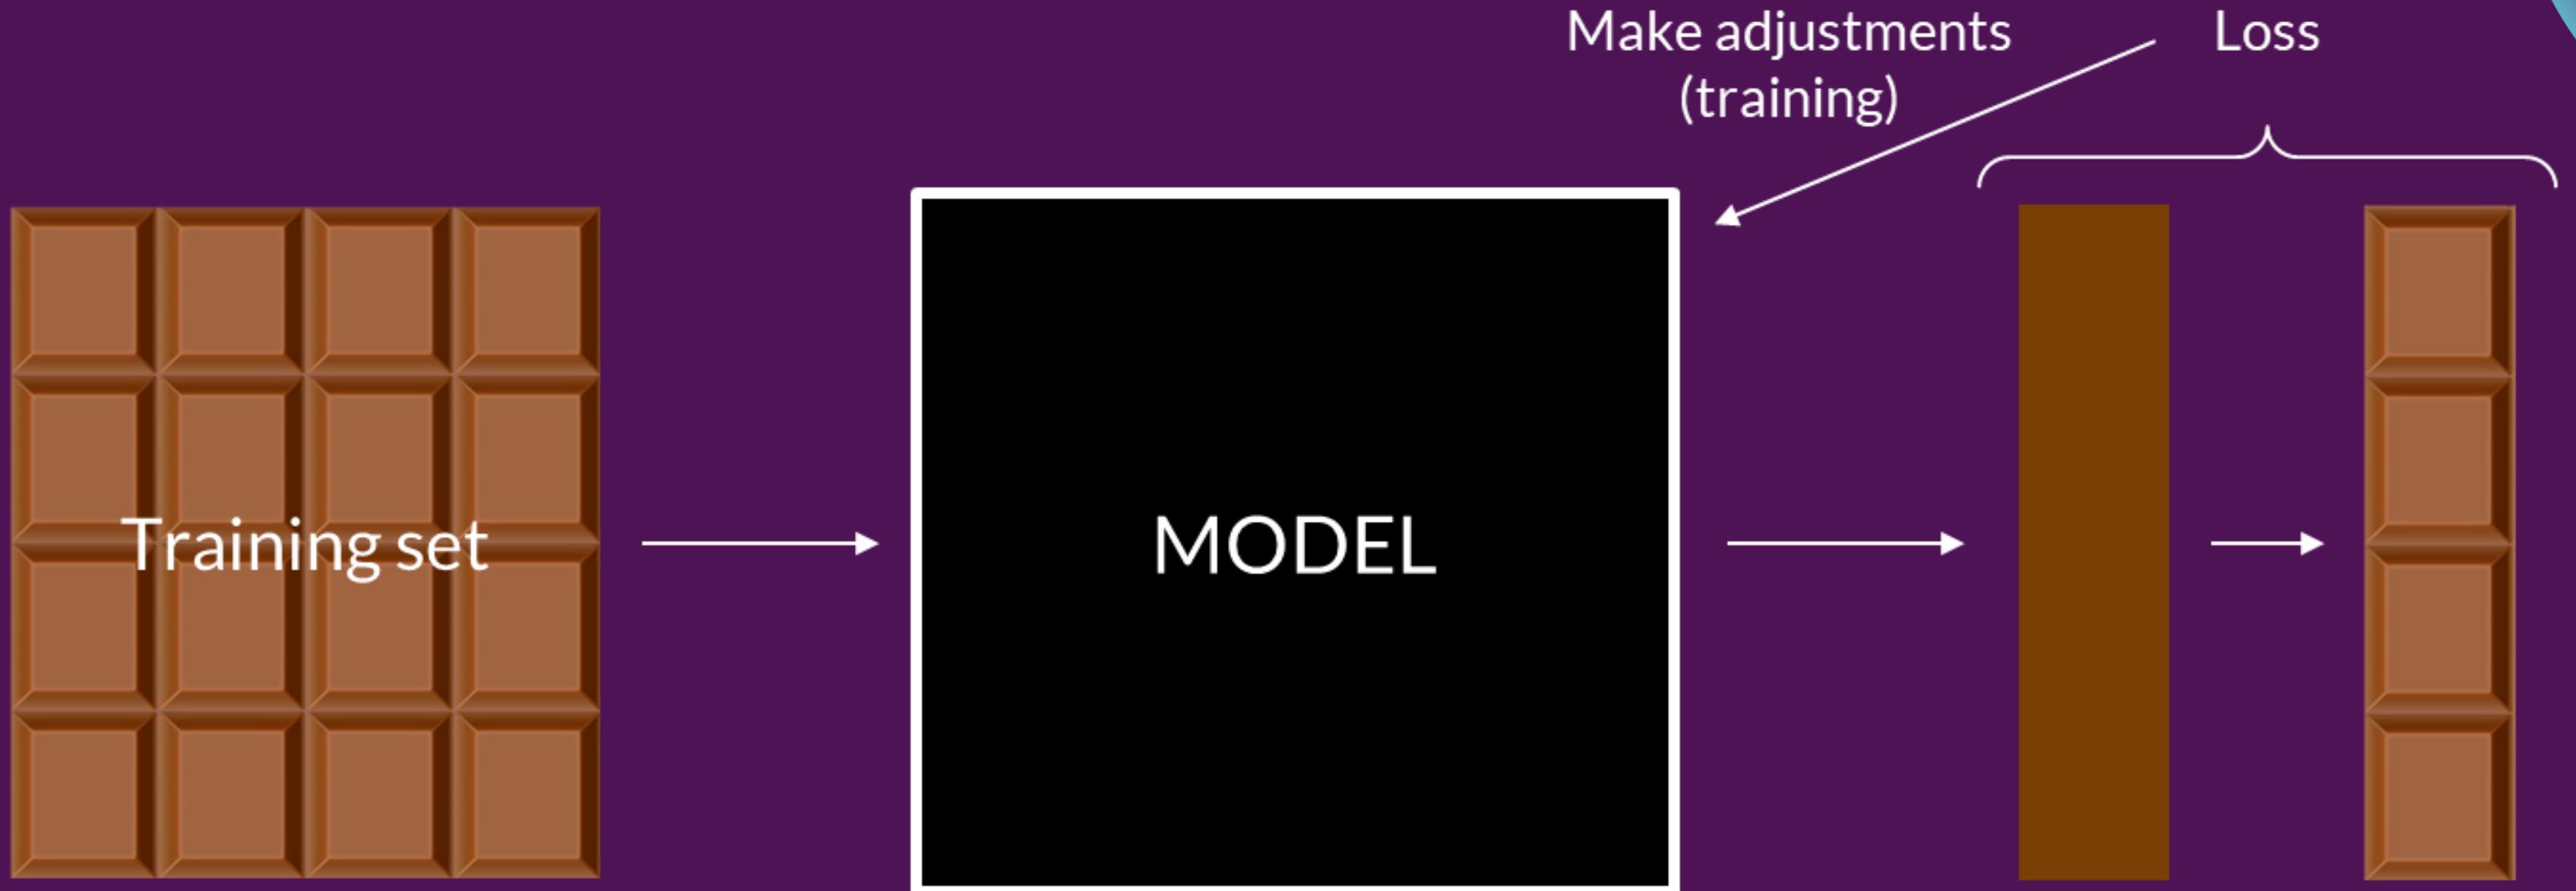
Validation dataset



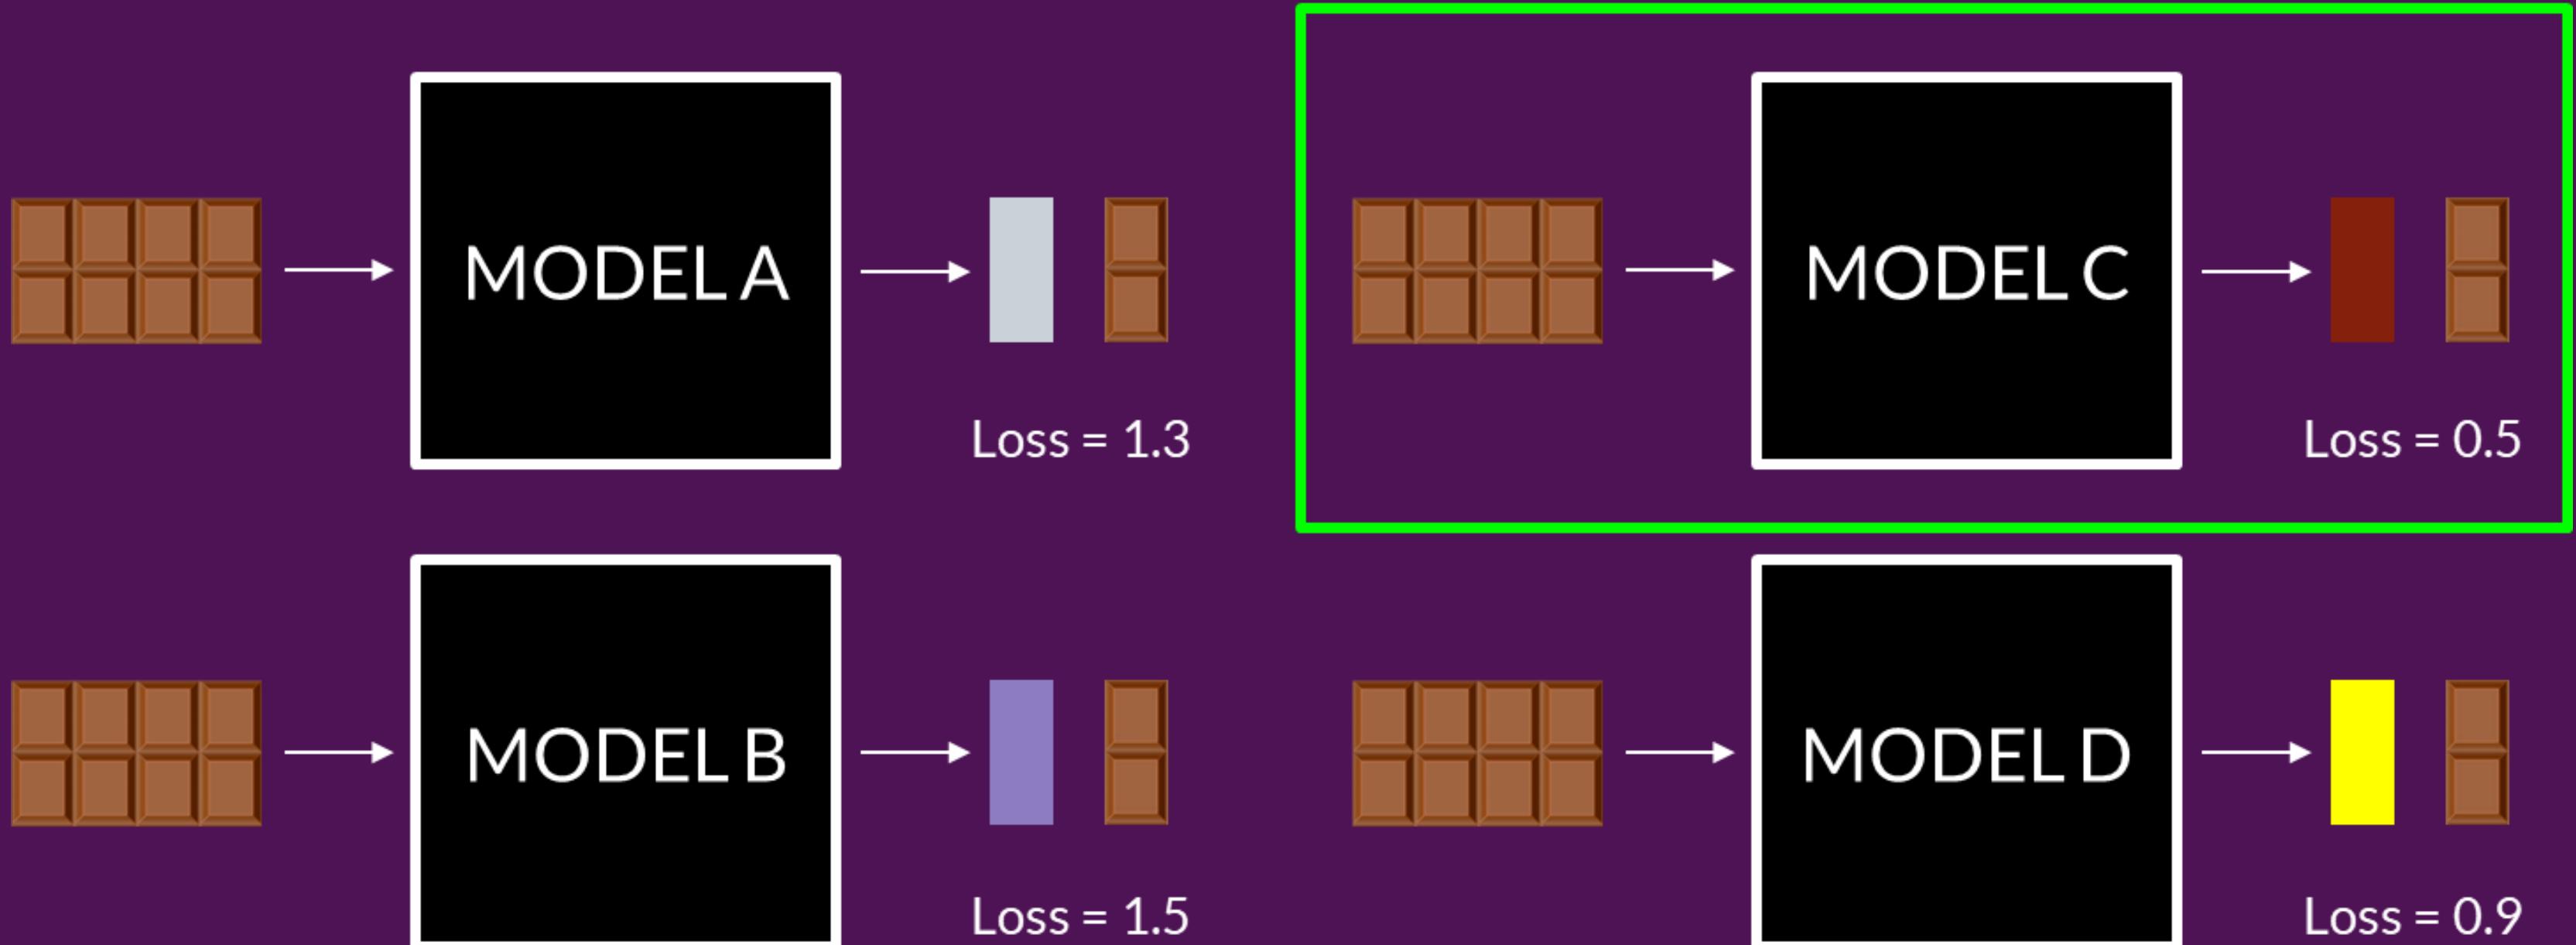
Testing dataset



# Supervised Learning Dataset



# Supervised Learning Dataset



# Metrics of Performance

## Classification Tasks

**True Positives (TP):** The model predicted the positive class correctly, and the actual class is indeed positive.

**True Negatives (TN):** The model predicted the negative class correctly, and the actual class is indeed negative.

**False Positives (FP):** The model predicted the positive class, but the actual class is negative.

**False Negatives (FN):** The model predicted the negative class, but the actual class is positive.



# Metrics of Performance

## Classification Tasks

**Accuracy :** Is the metric to evaluate how correct the model is to predict the values.

$$\text{Accuracy} = (\text{TP} + \text{TN}) / (\text{TP} + \text{FP} + \text{FN} + \text{TN})$$

**Precision :** Out of all actual values, how many are correctly predicted

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

**Recall :** From all the predicted values how many are correctly predicted.

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

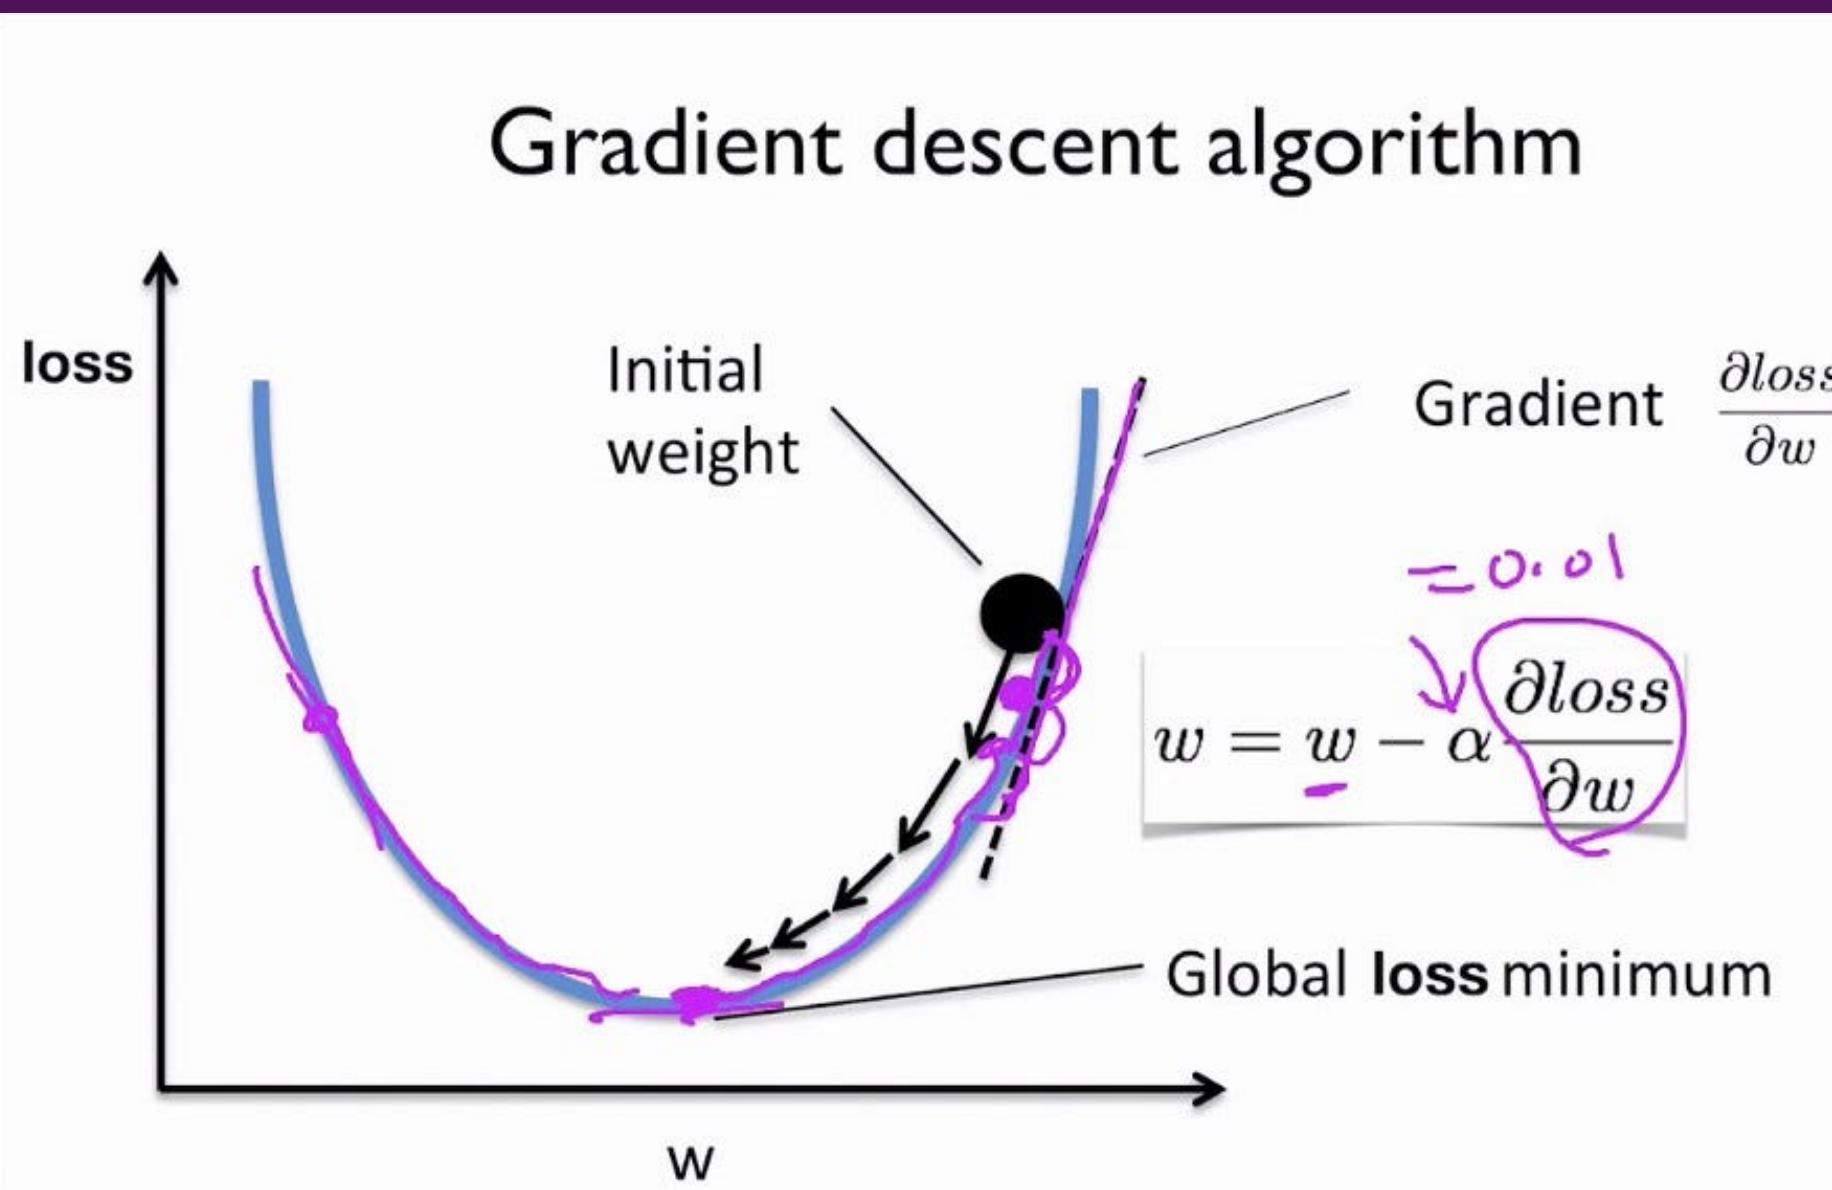
**F1 score :** Is the best choice if your distribution is uneven.

$$\text{F1\_score} = 2(\text{Recall Precision}) / (\text{Recall} + \text{Precision})$$



# Optimization Techniques

Optimization refers to the process of adjusting the parameters of a model to minimize or maximize some objective function.



Gradient Descent

$$\nu_t = \beta_1 * \nu_{t-1} - (1 - \beta_1) * g_t$$

$$s_t = \beta_2 * s_{t-1} - (1 - \beta_2) * g_t^2$$

$$\Delta \omega_t = -\eta \frac{\nu_t}{\sqrt{s_t + \epsilon}} * g_t$$

$$\omega_{t+1} = \omega_t + \Delta \omega_t$$

$\eta$  : Initial Learning rate

$g_t$  : Gradient at time  $t$  along  $\omega^j$

$\nu_t$  : Exponential Average of gradients along  $\omega_j$

$s_t$  : Exponential Average of squares of gradients along  $\omega_j$

$\beta_1, \beta_2$  : Hyperparameters

ADAM = momentum + Root Mean Square Propagation

# Regularization Techniques

Regularization is a set of techniques used to prevent overfitting and improve the generalization ability of a model.

L2 Regularization (Weight Decay) : adds a penalty term to the loss function, which penalizes large weights in the network.

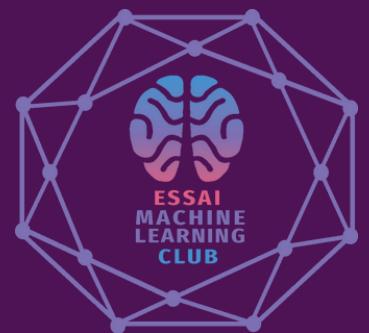
Dropout : technique where randomly selected neurons are ignored during training.

Data Augmentation : involves artificially increasing the size of the training dataset by applying transformations such as rotation, translation, scaling, flipping, and cropping to the input images.

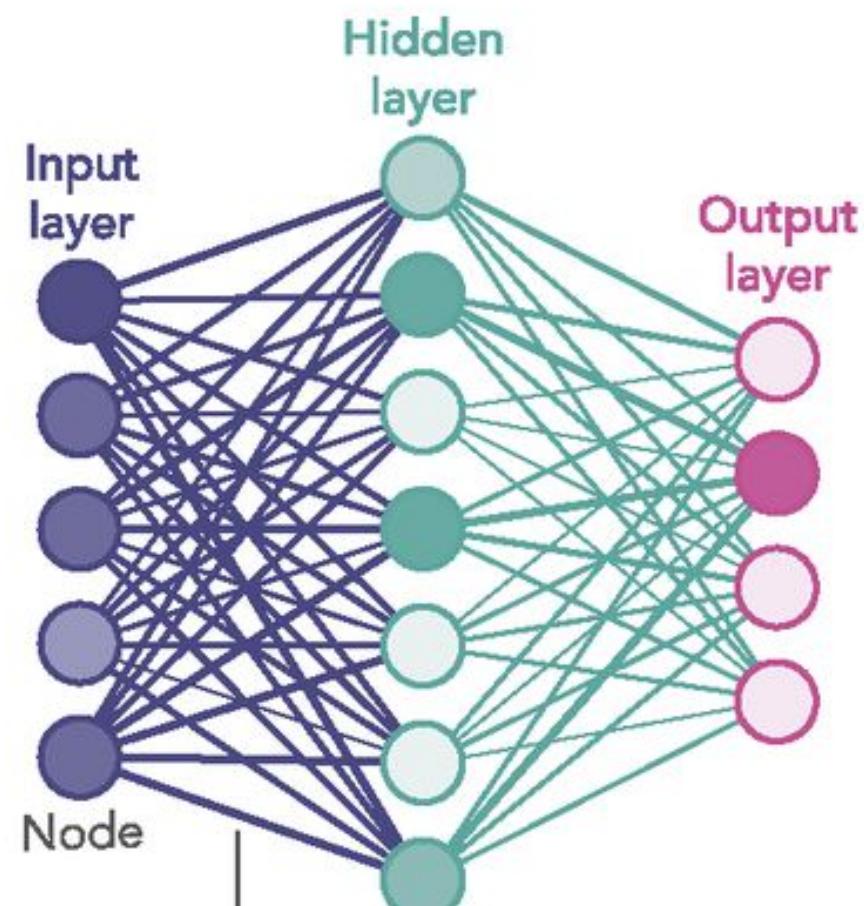
Batch Normalization : stabilizing training by reducing internal covariate shift and acting as a regularizer by adding noise.

Early Stopping : stopping the training process when the performance starts to degrade.

# NEURAL NETS



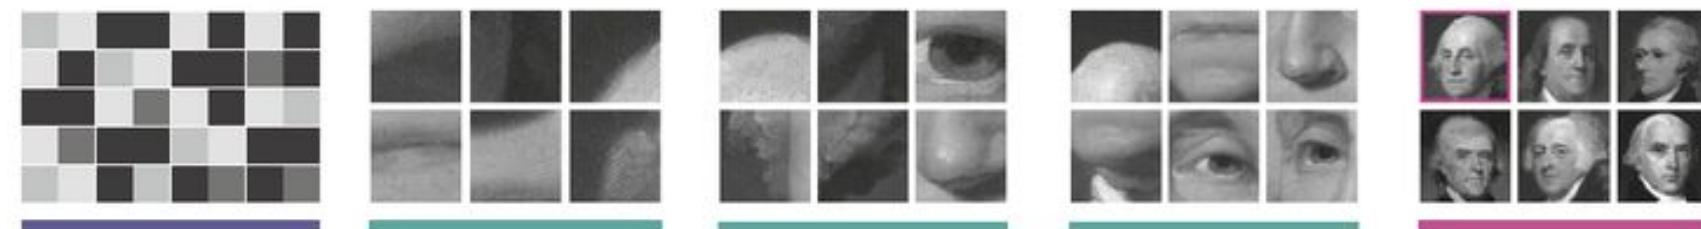
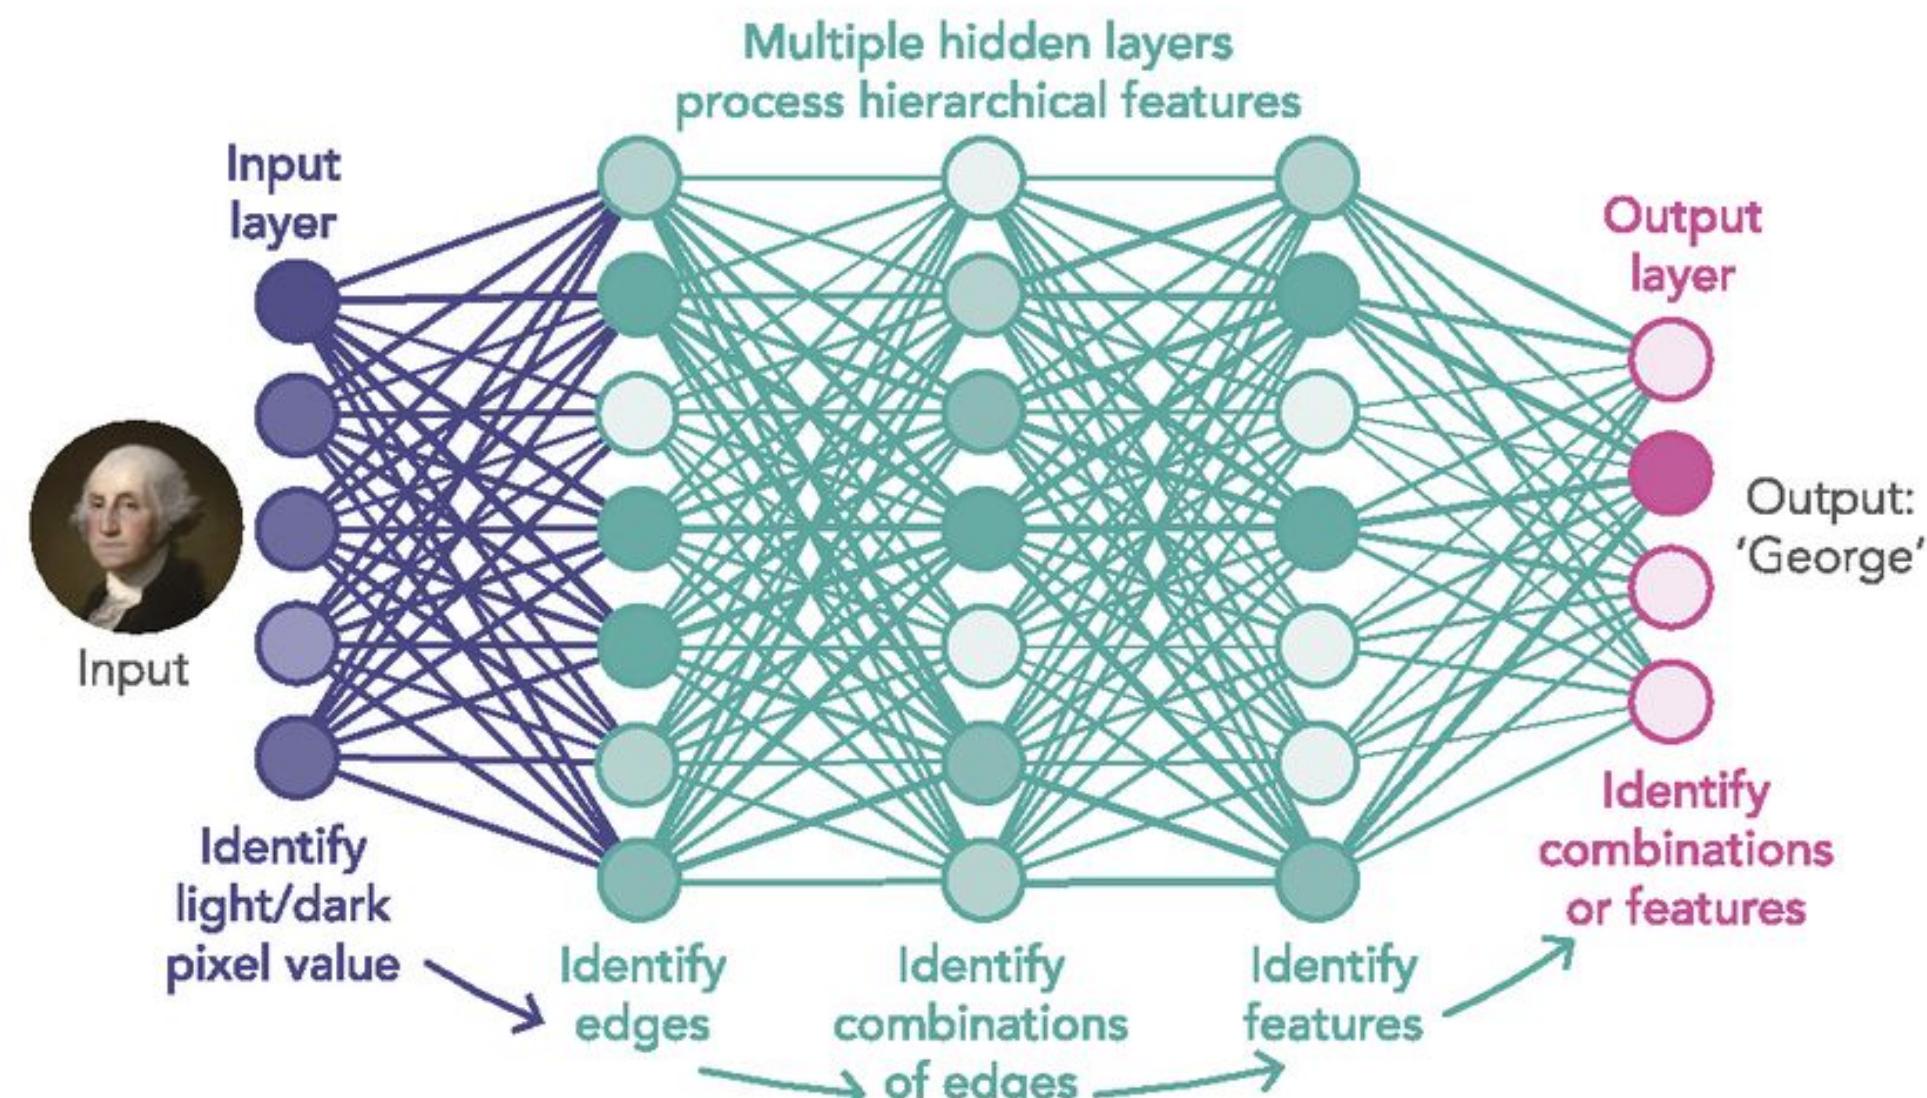
## 1980S-ERA NEURAL NETWORK



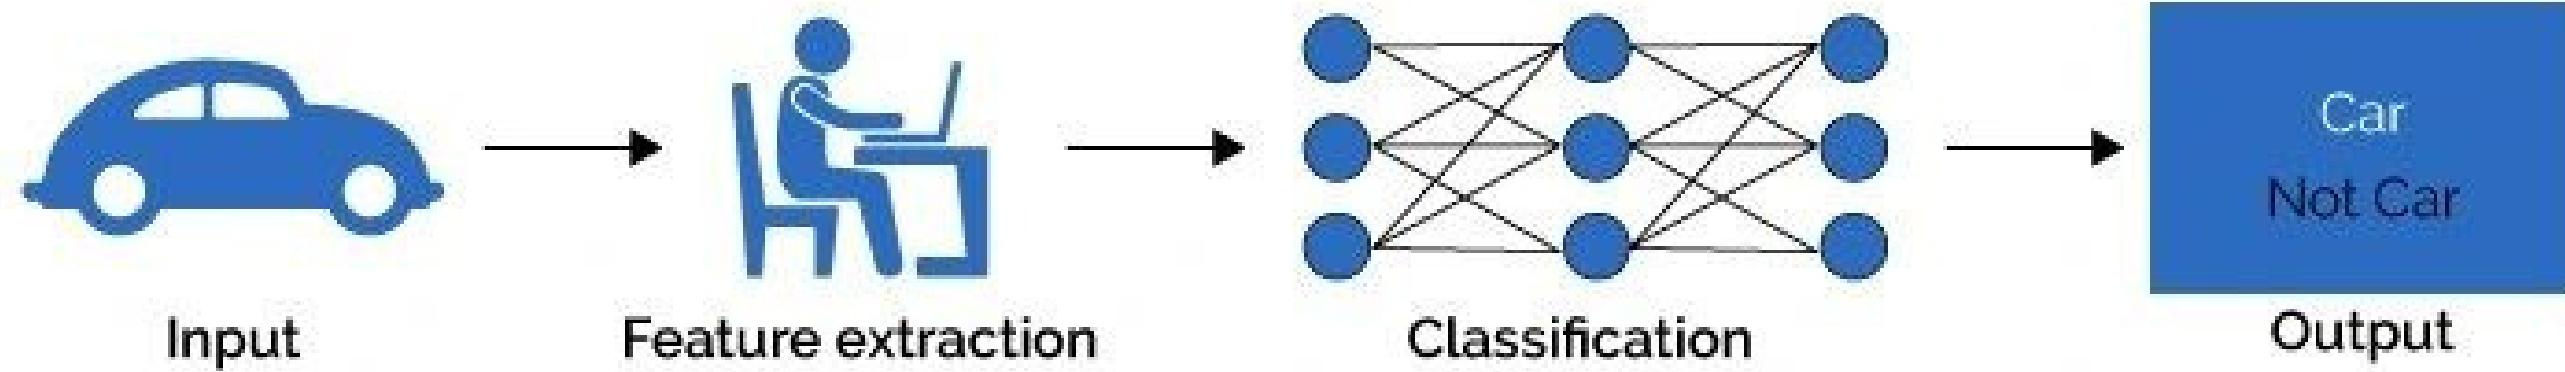
Links carry signals from one node to another, boosting or damping them according to each link's 'weight'.



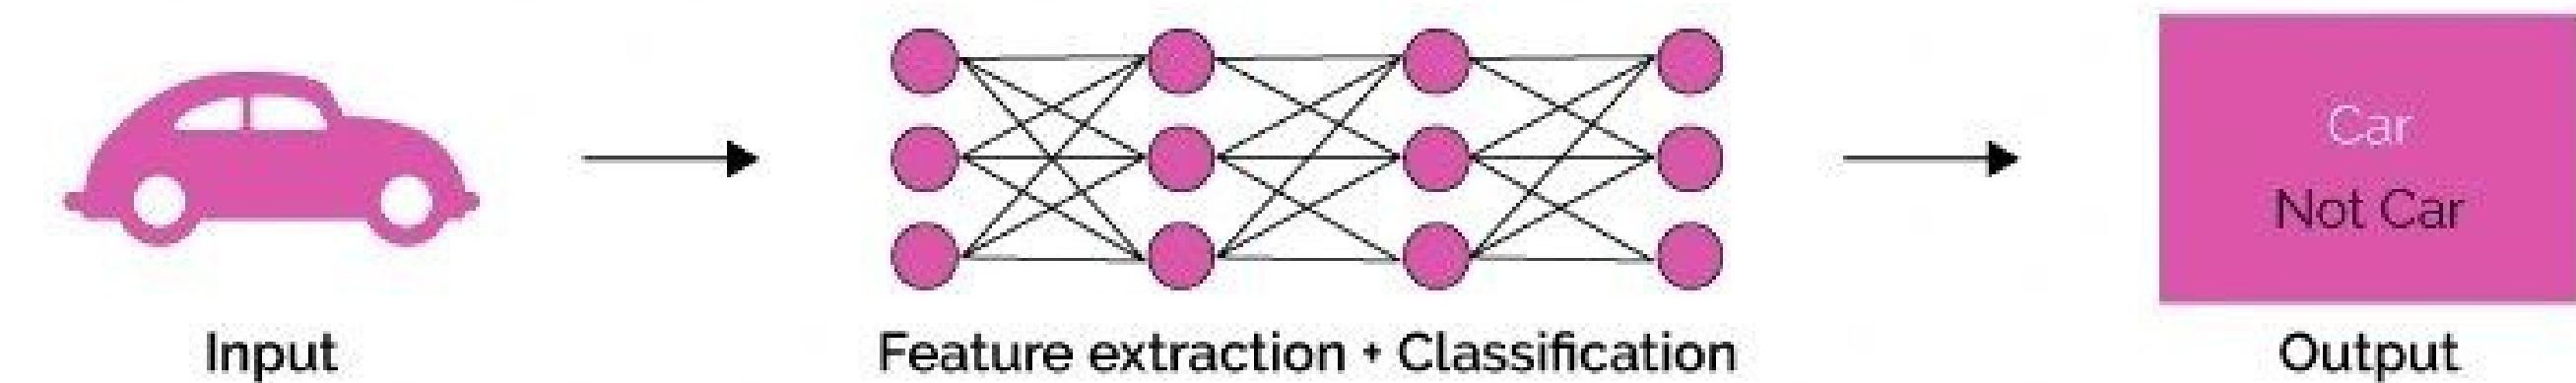
## DEEP LEARNING NEURAL NETWORK



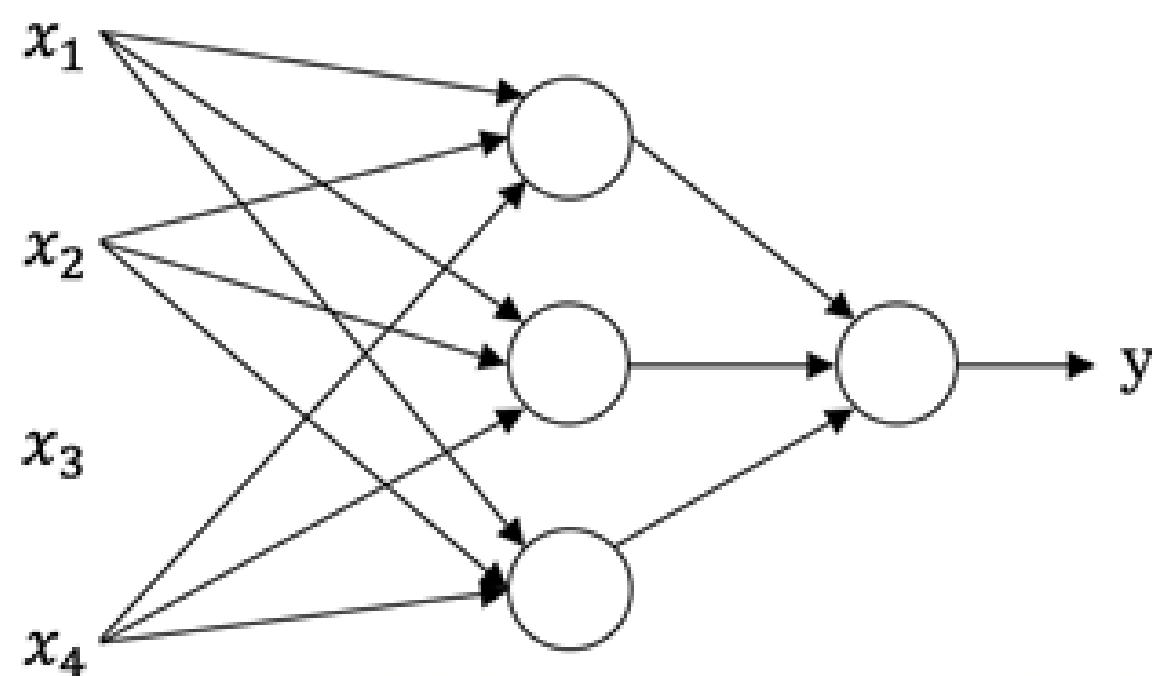
## Machine Learning



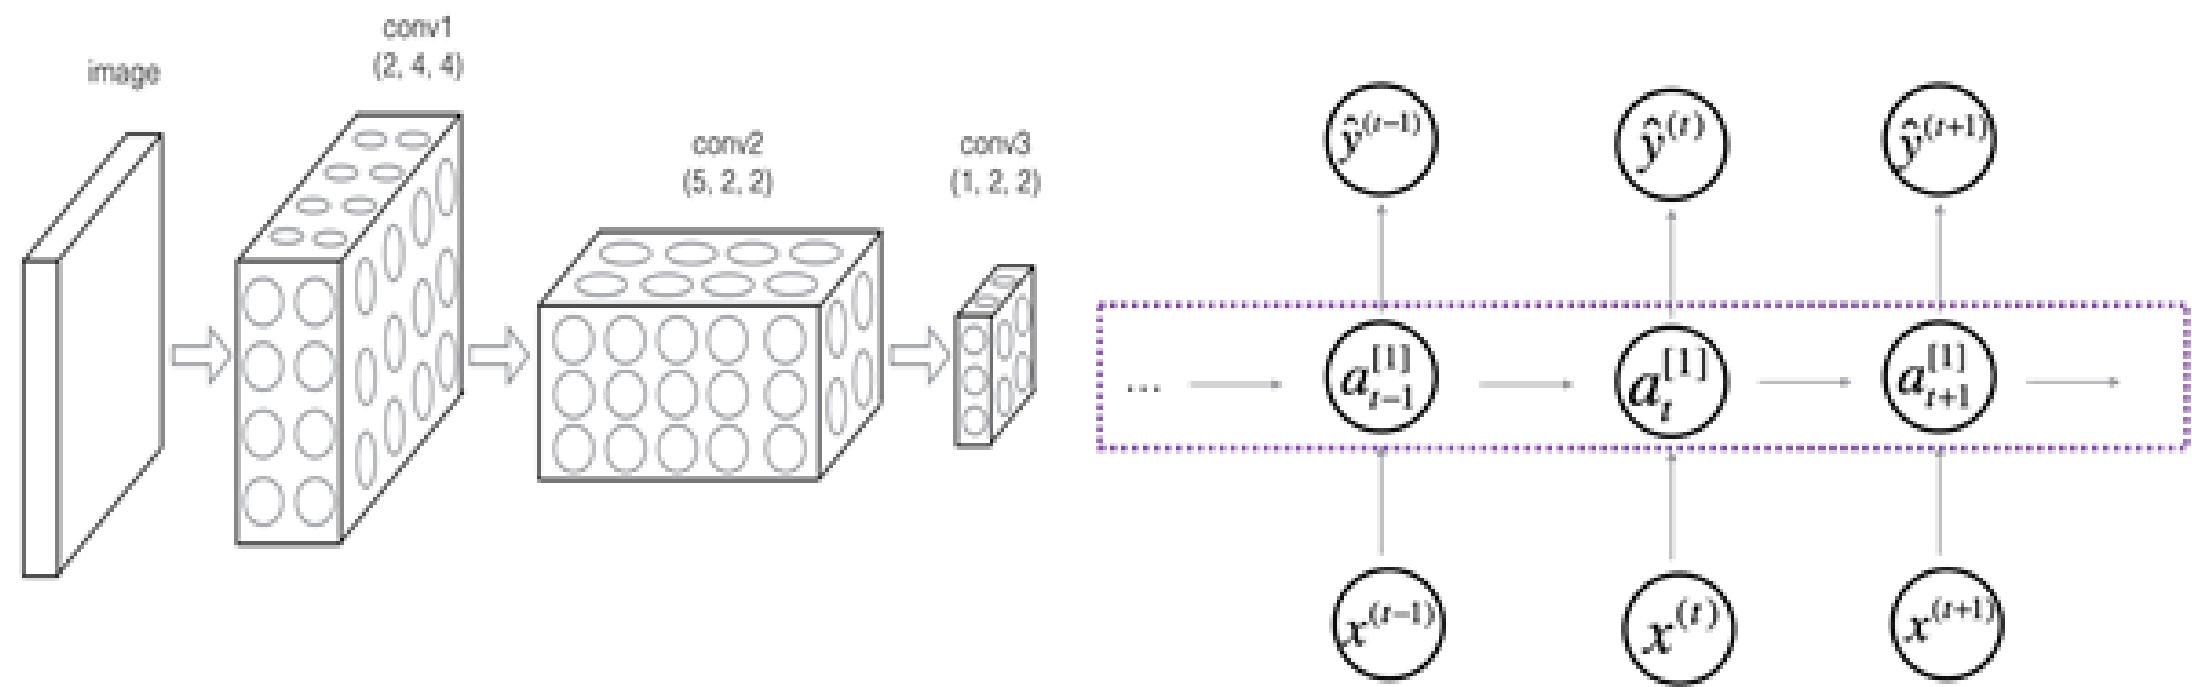
## Deep Learning



# Neural Network examples



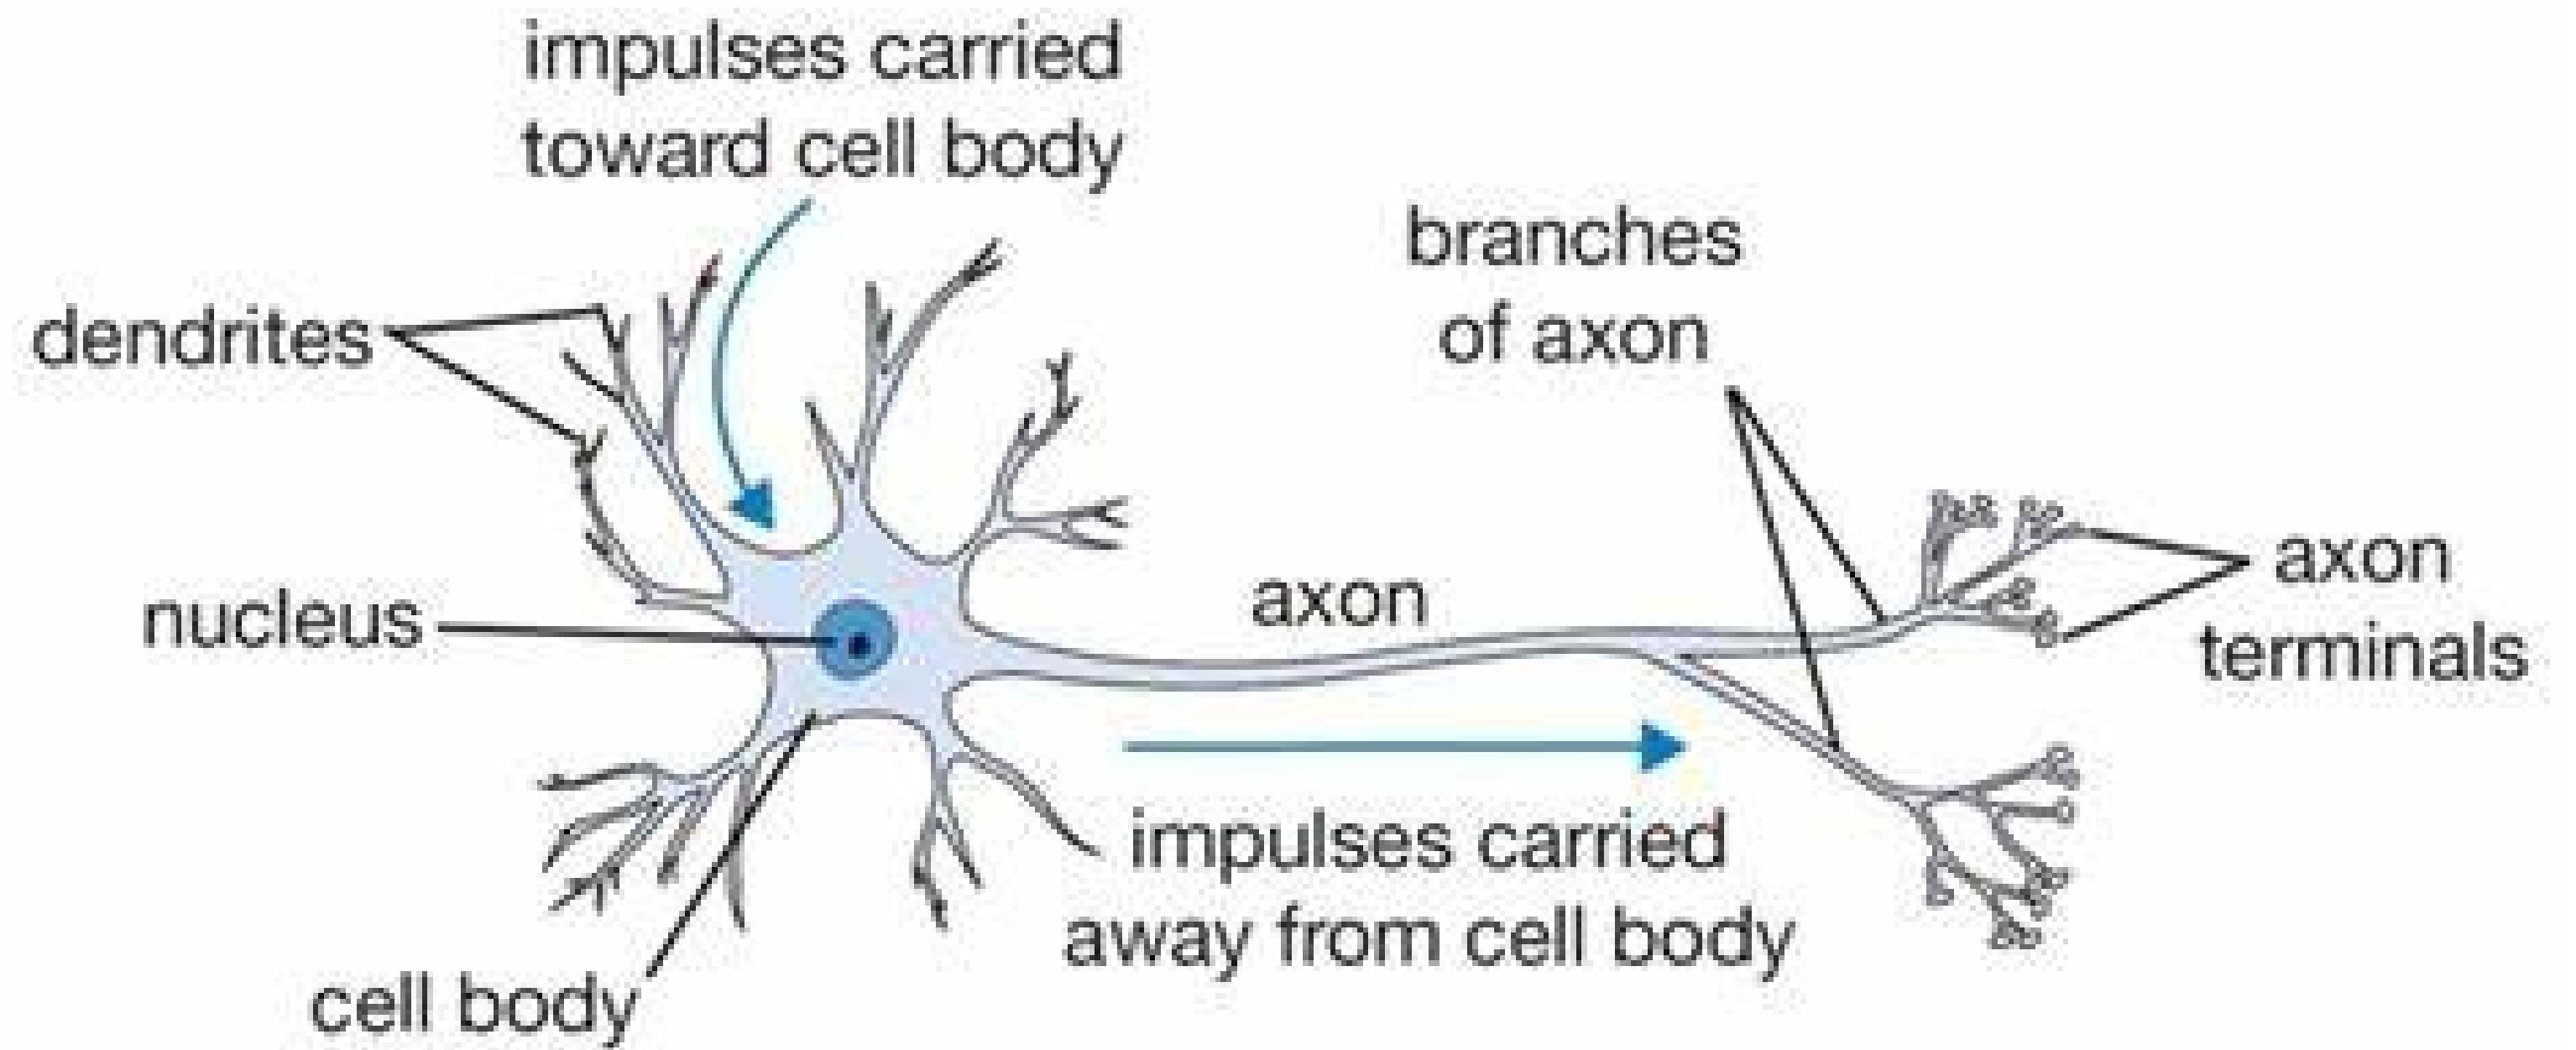
Standard NN

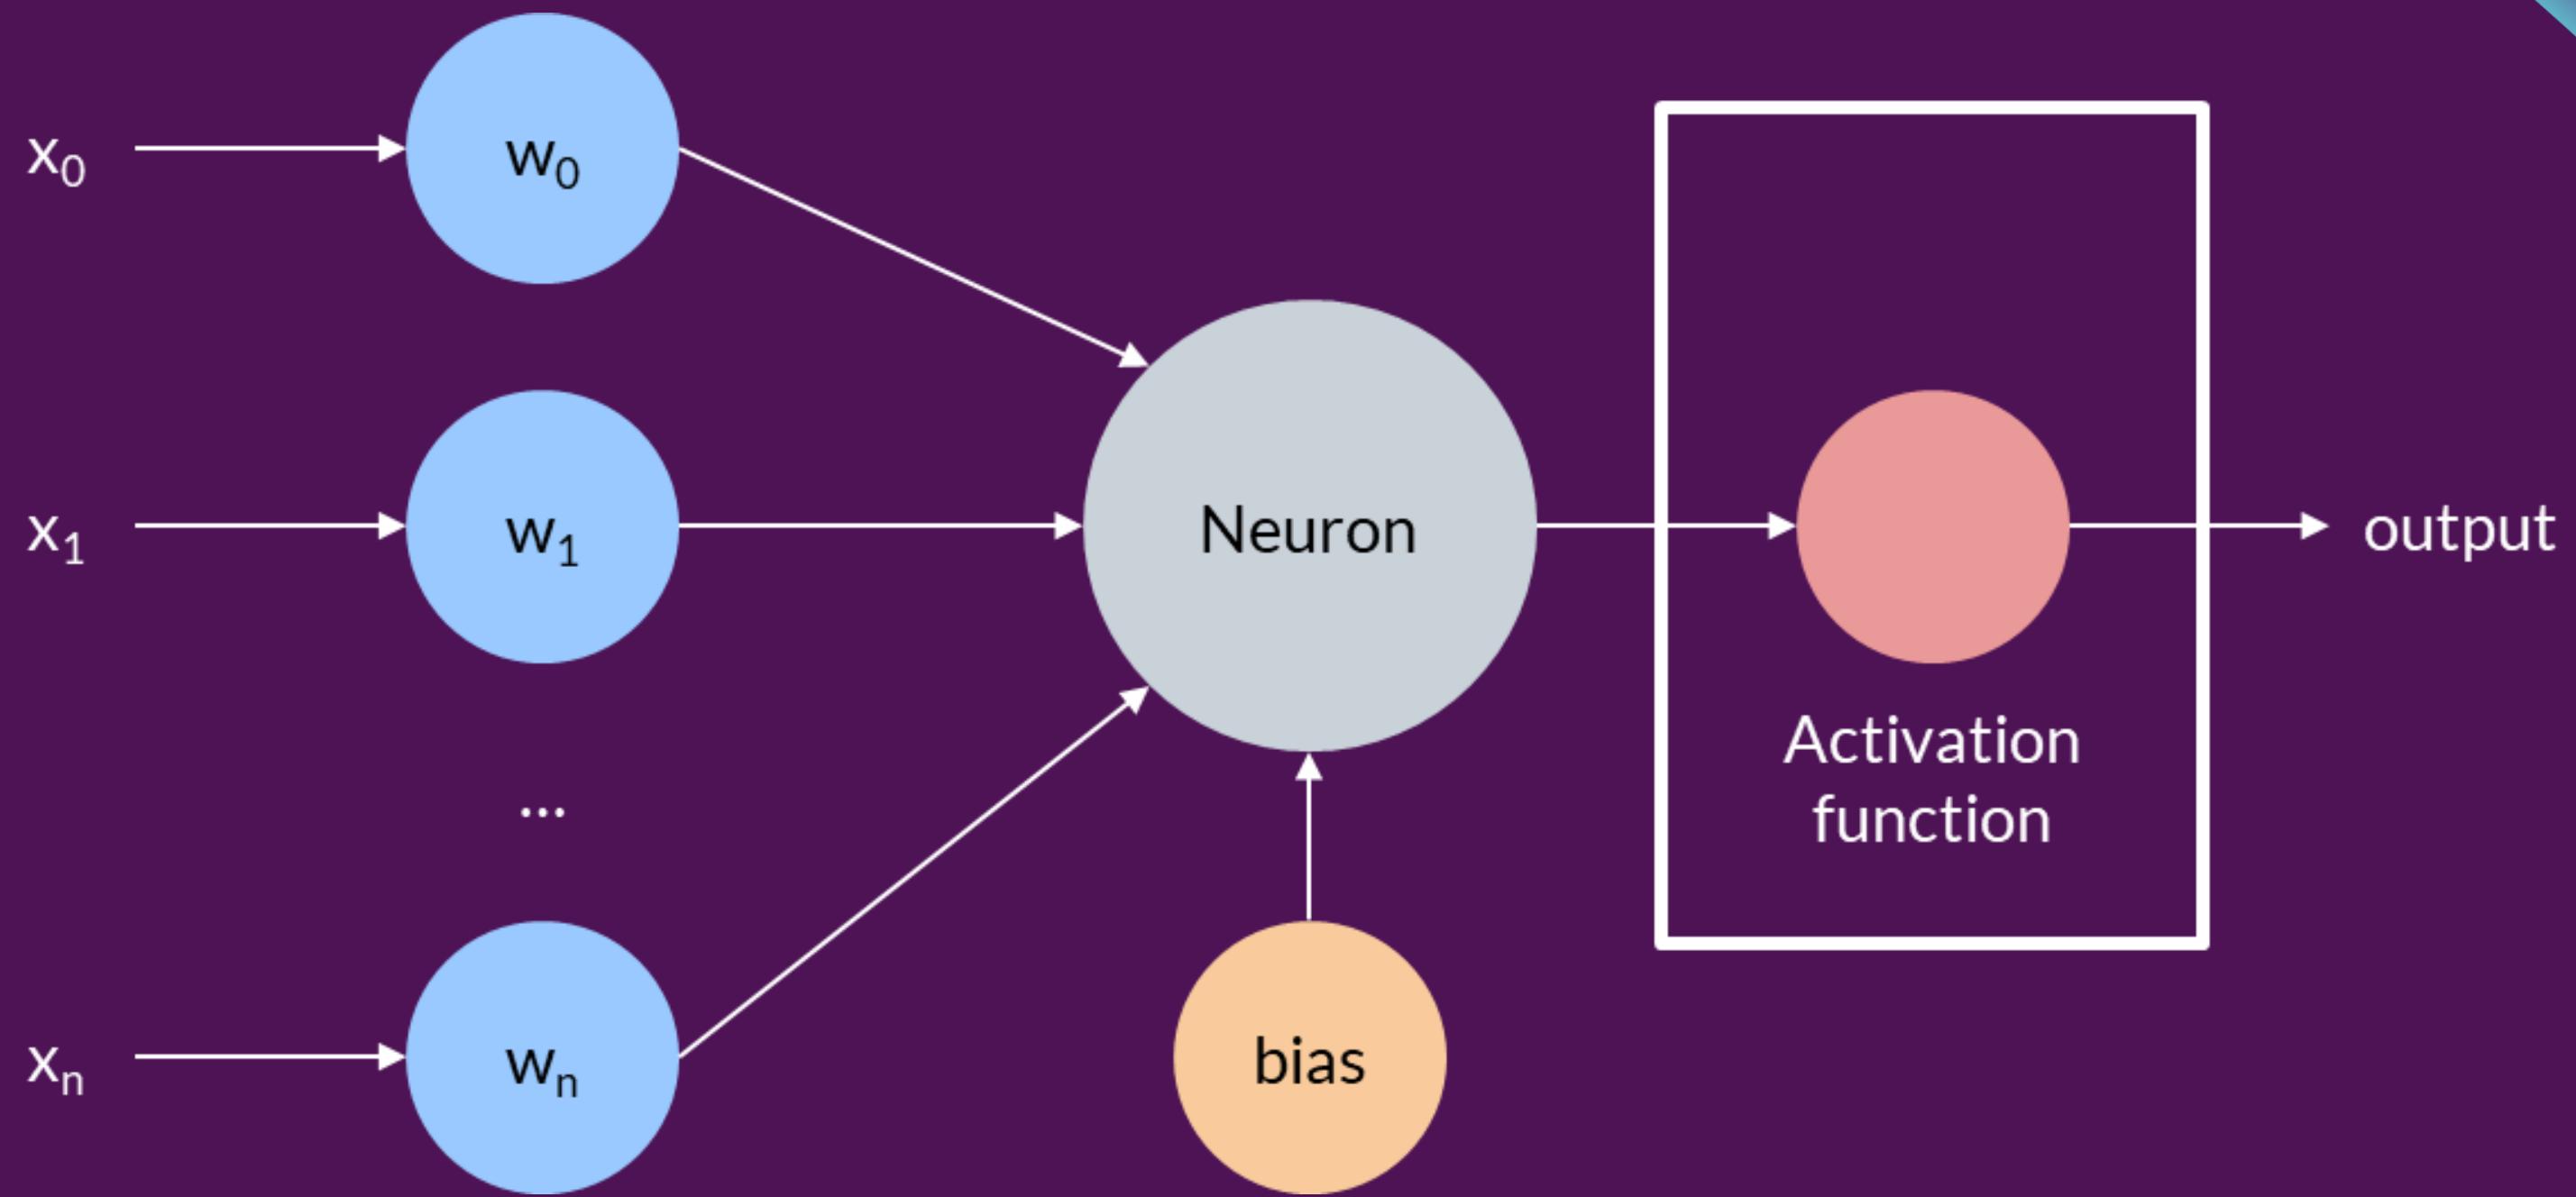


Convolutional NN

Recurrent NN

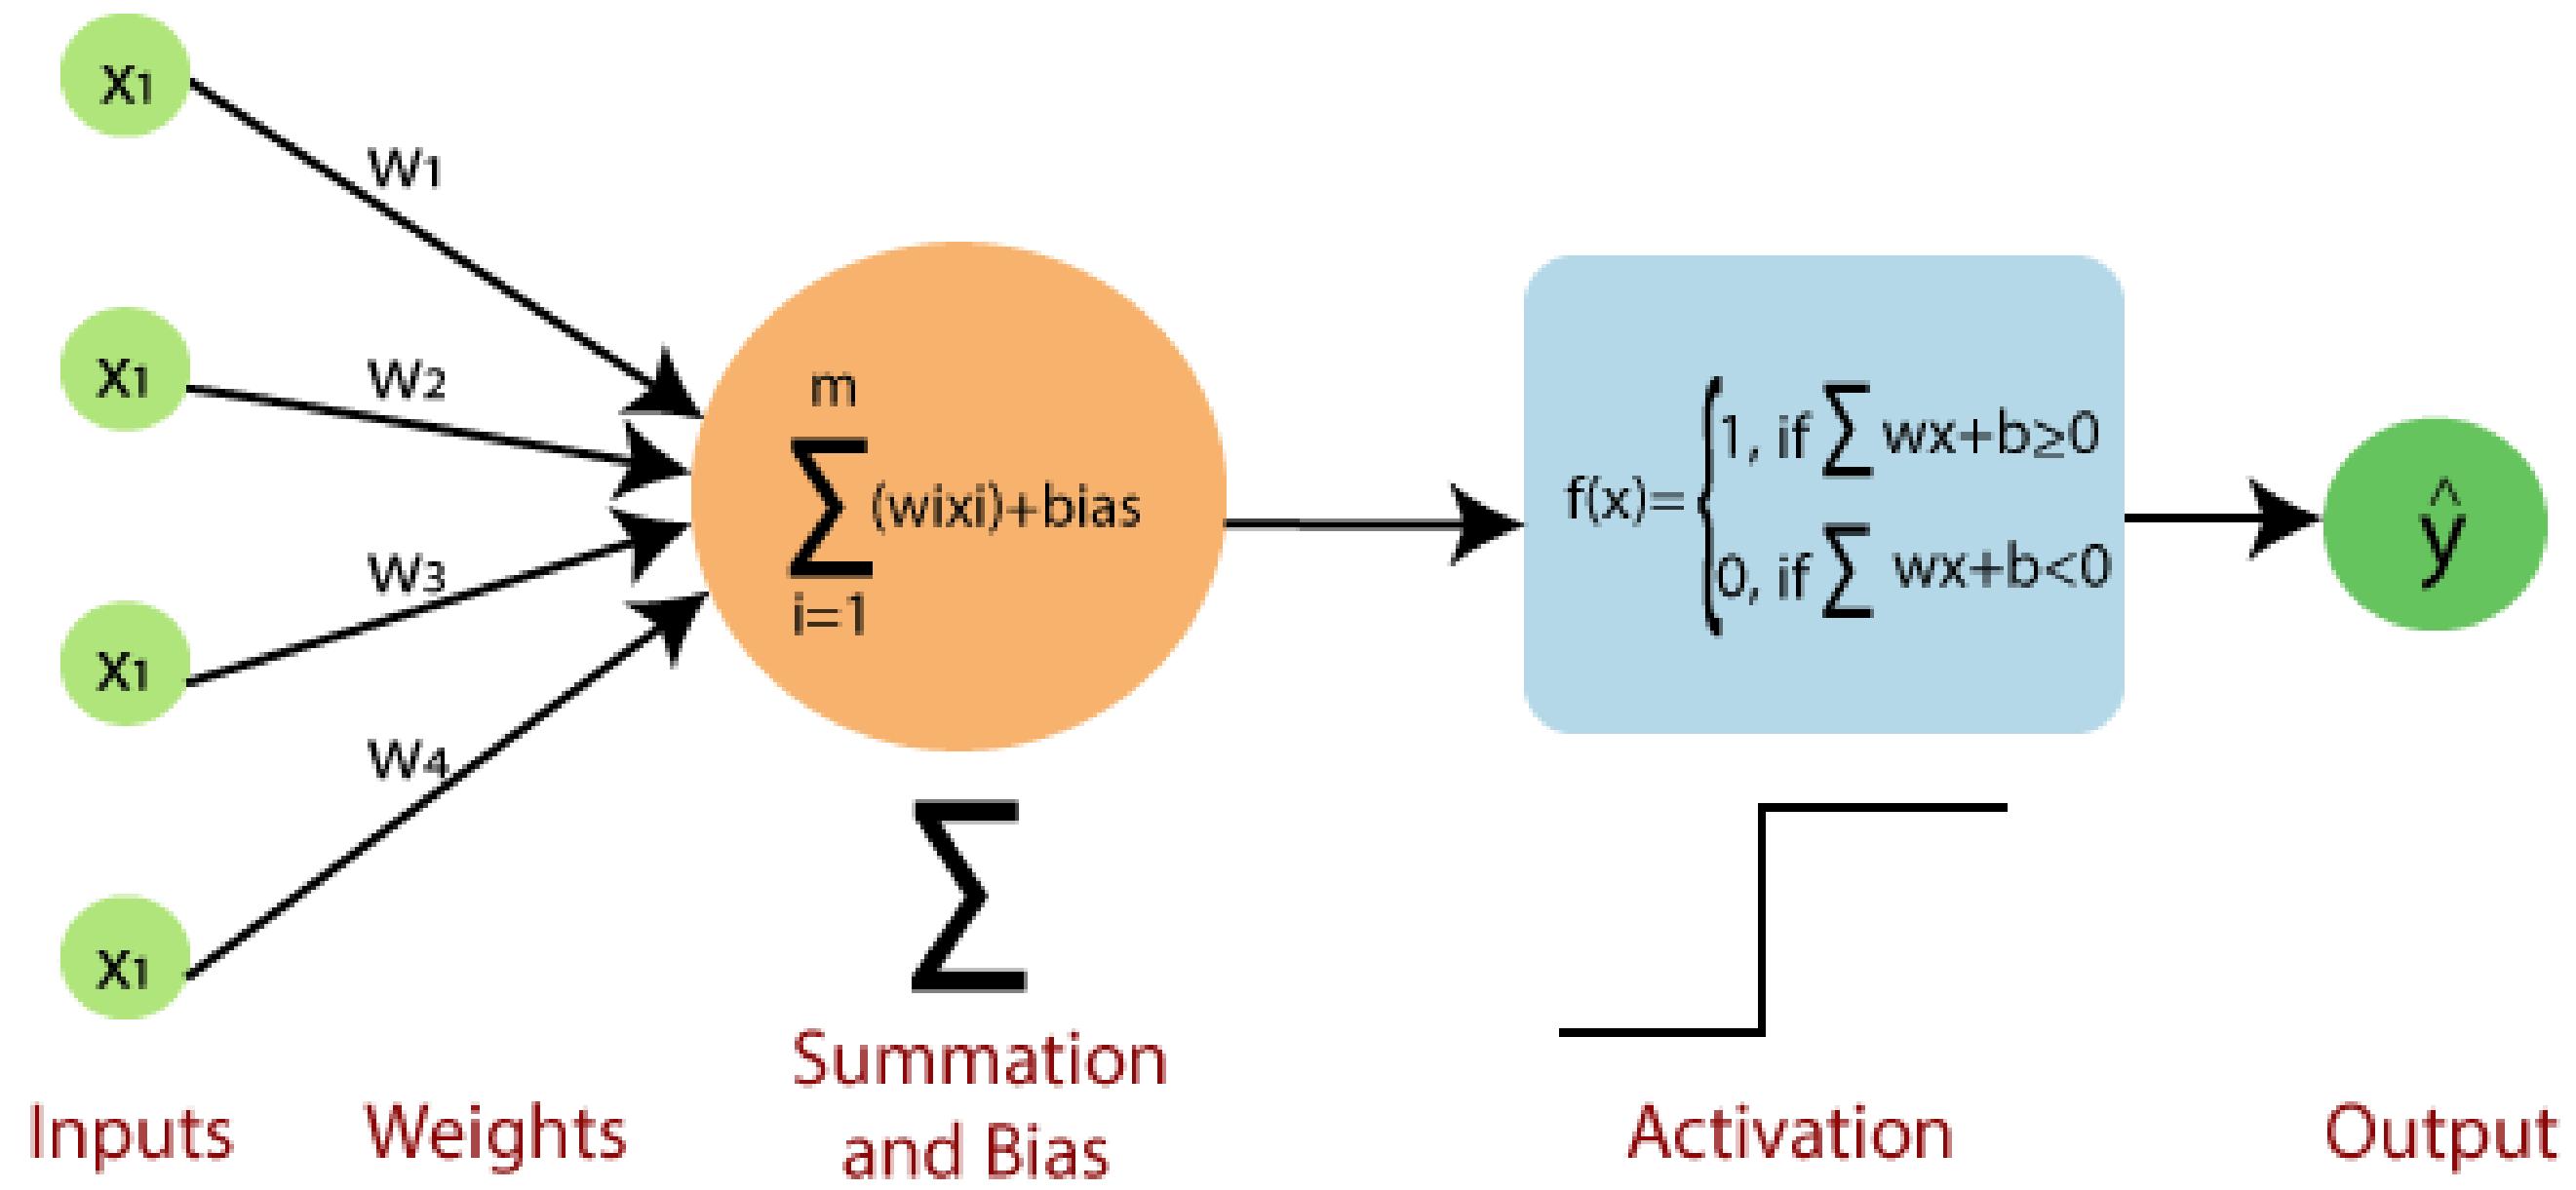






A node, also called a neuron, is a unit that contains a value calculated through a mathematical operation.

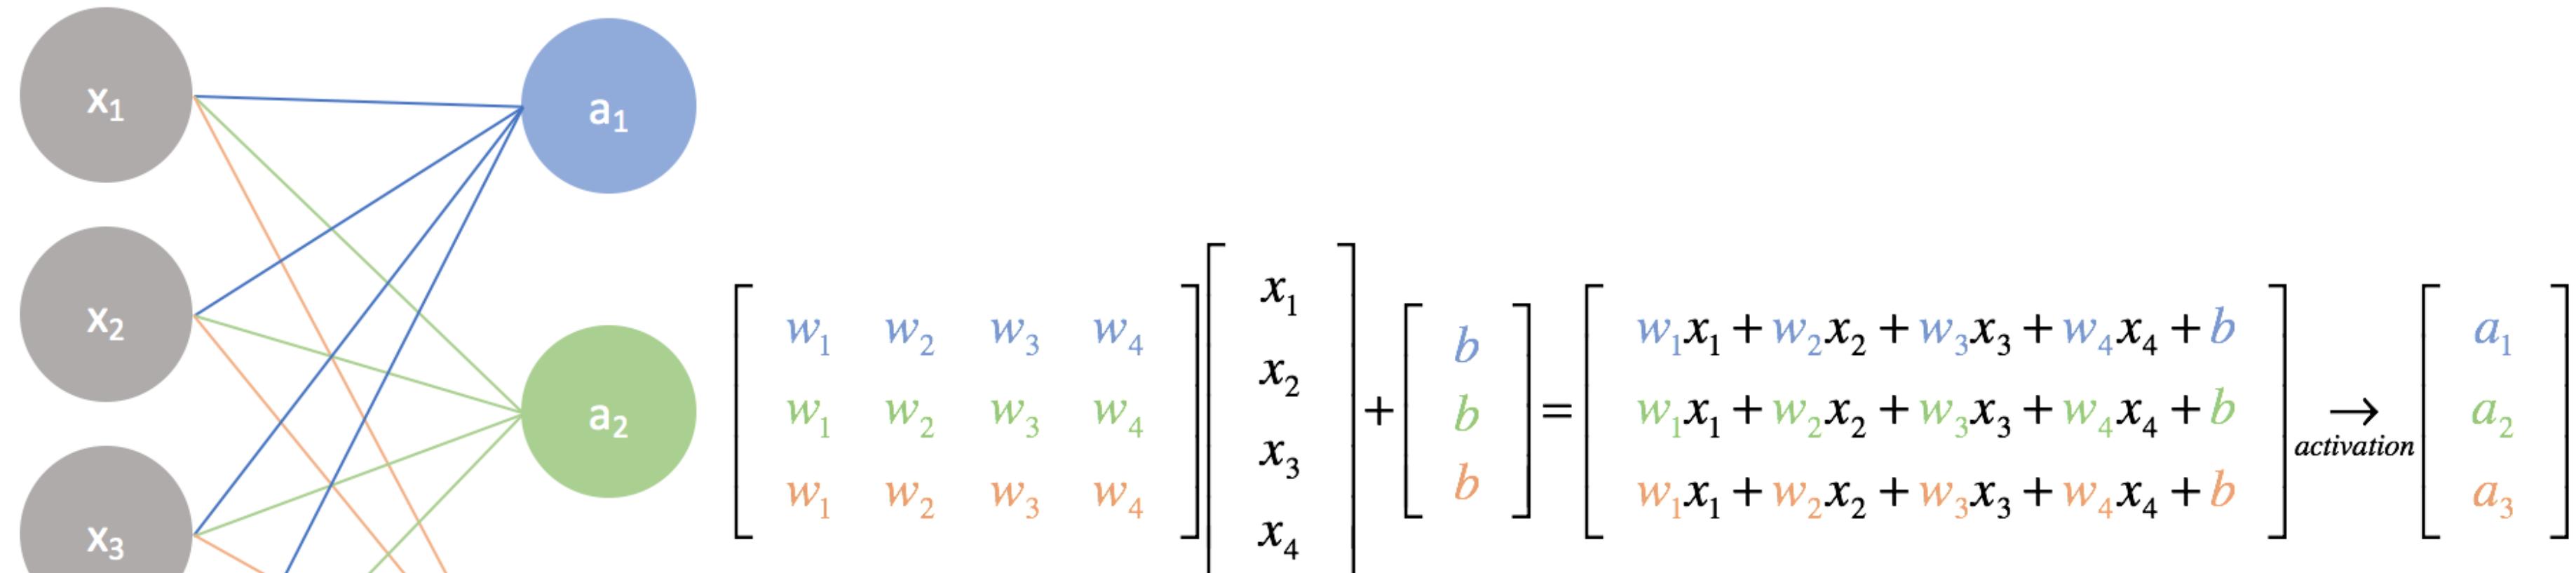




Input layer

Output layer

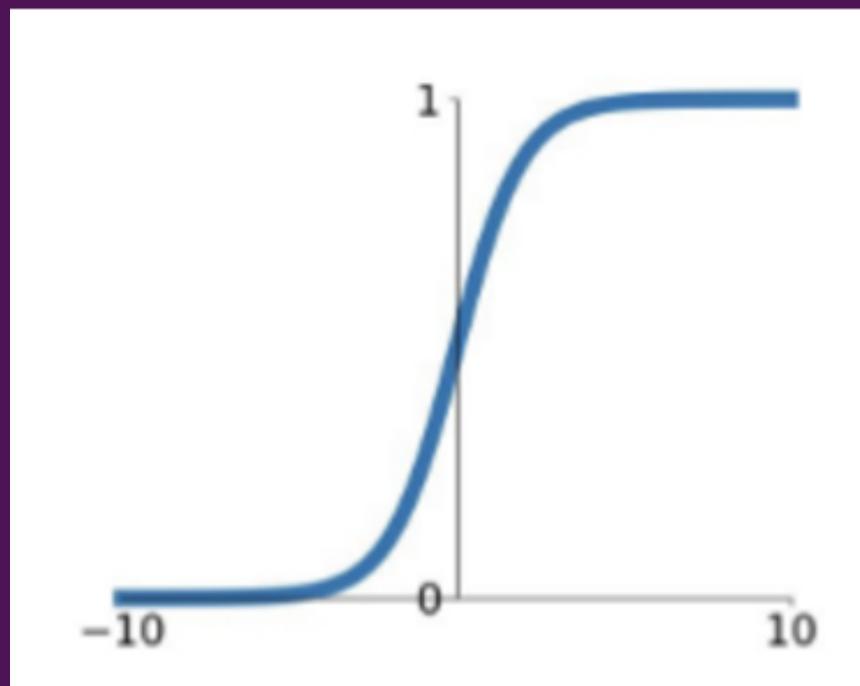
# A simple neural network



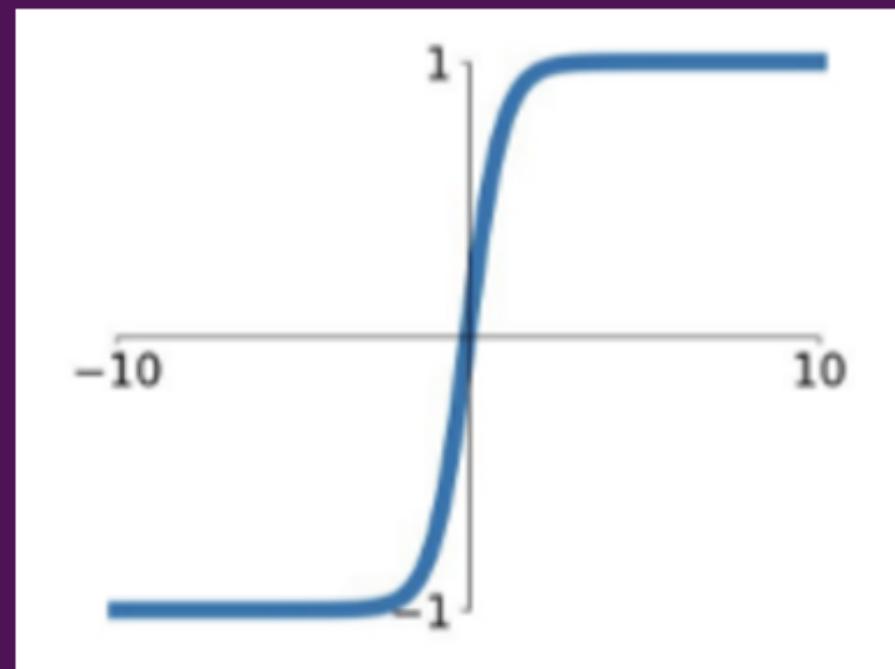
**Every connection is associated a weight (can be random at first)**  
**Every neuron is associated a bias (can be random at first)**

# Activation Fonctions

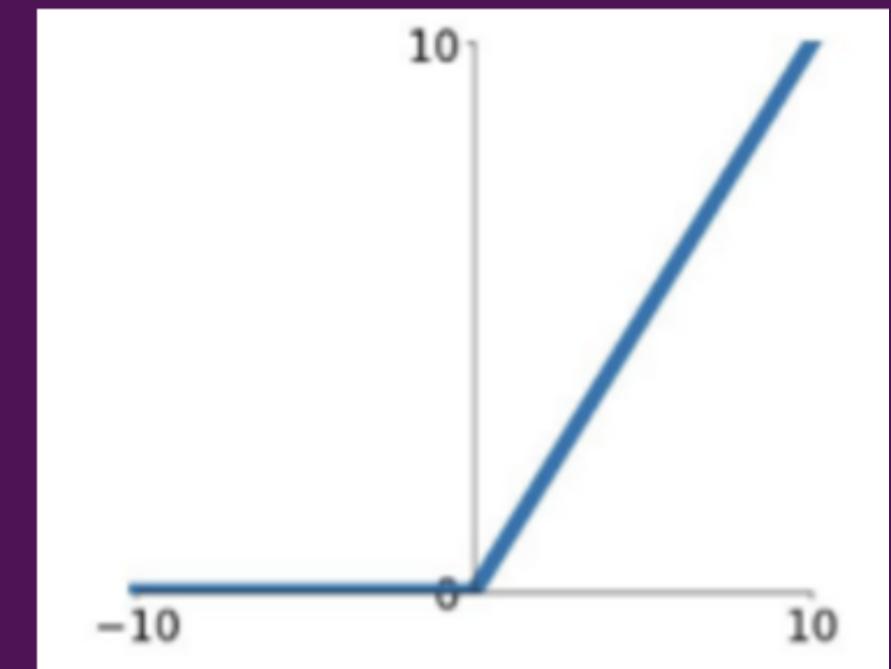
Sigmoid



Tanh



ReLU



$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

binary classification problem

$$f(x) = \frac{(e^x - e^{-x})}{(e^x + e^{-x})}$$

Descent is faster

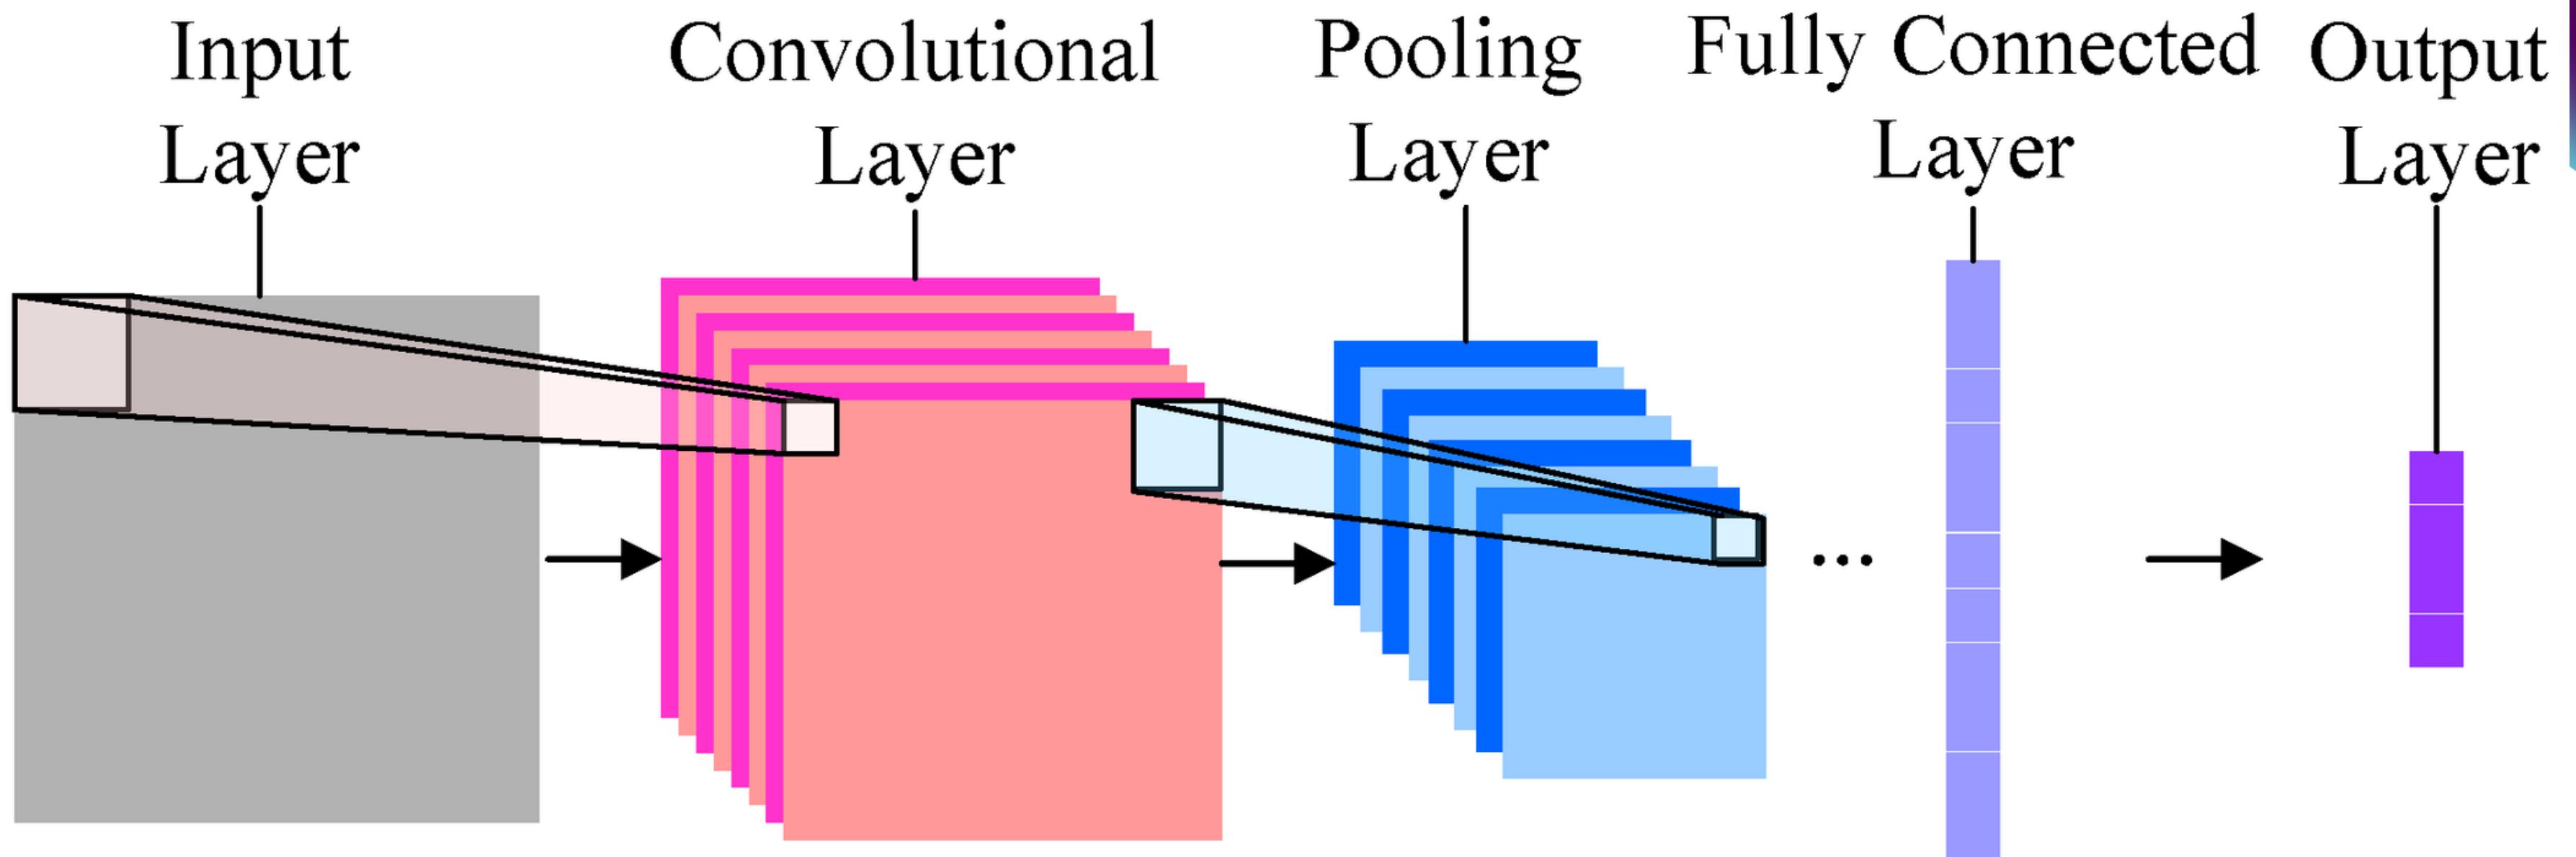
$$relu(x) = \max(0, x)$$

Default choice for activation



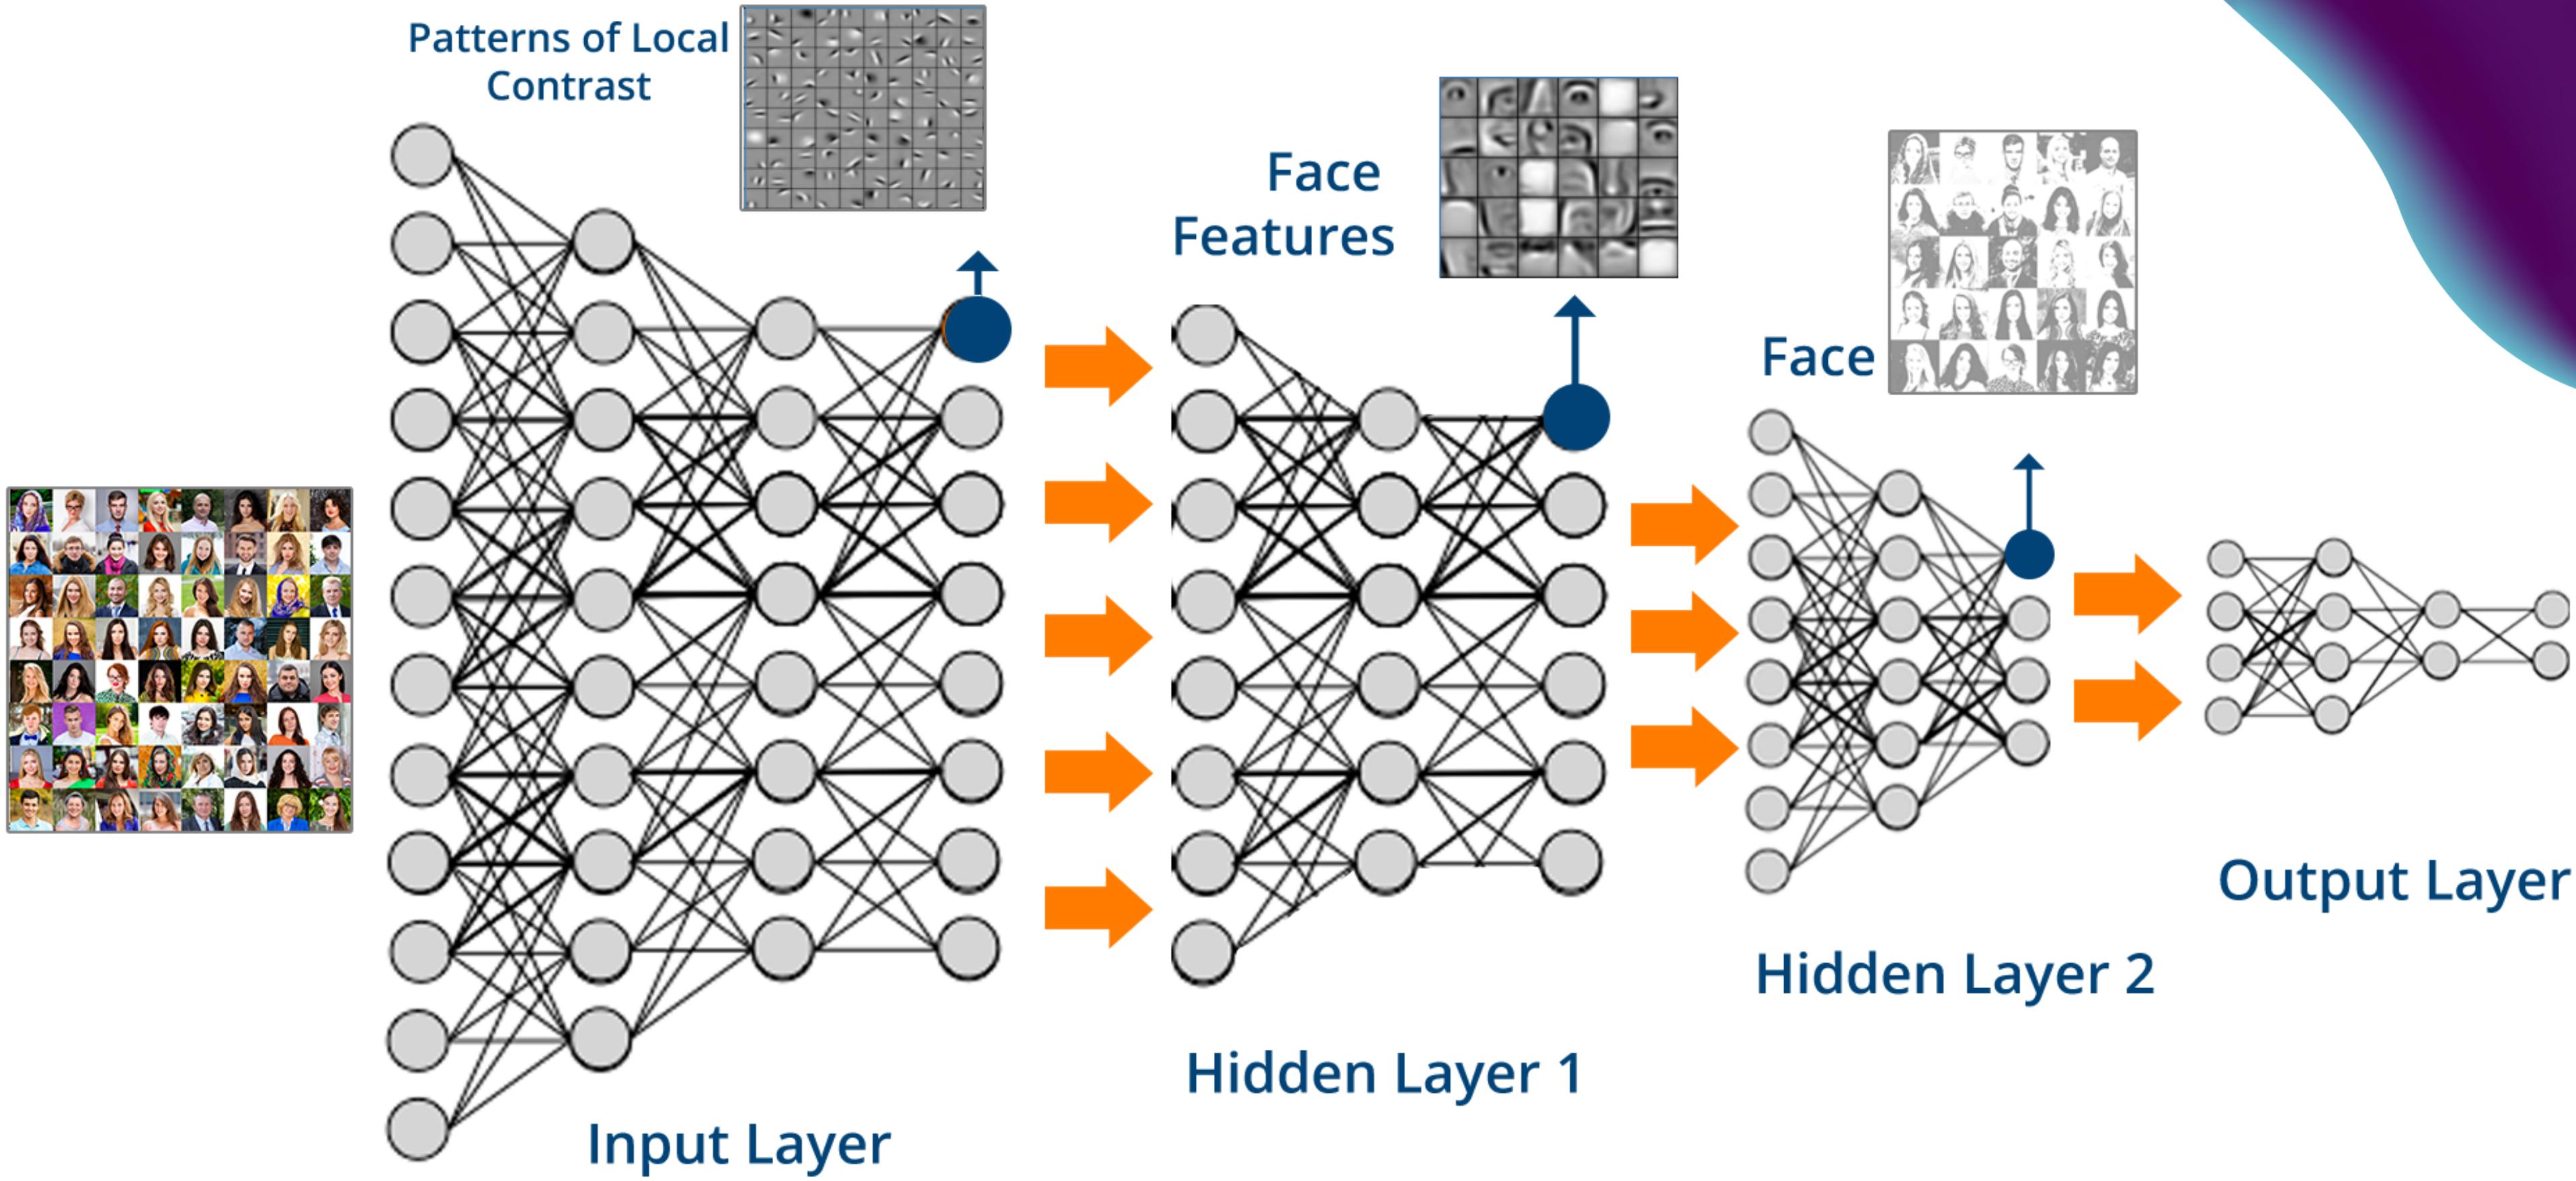
Without activation functions, this becomes a linear model!

# Convolutional Neural Network

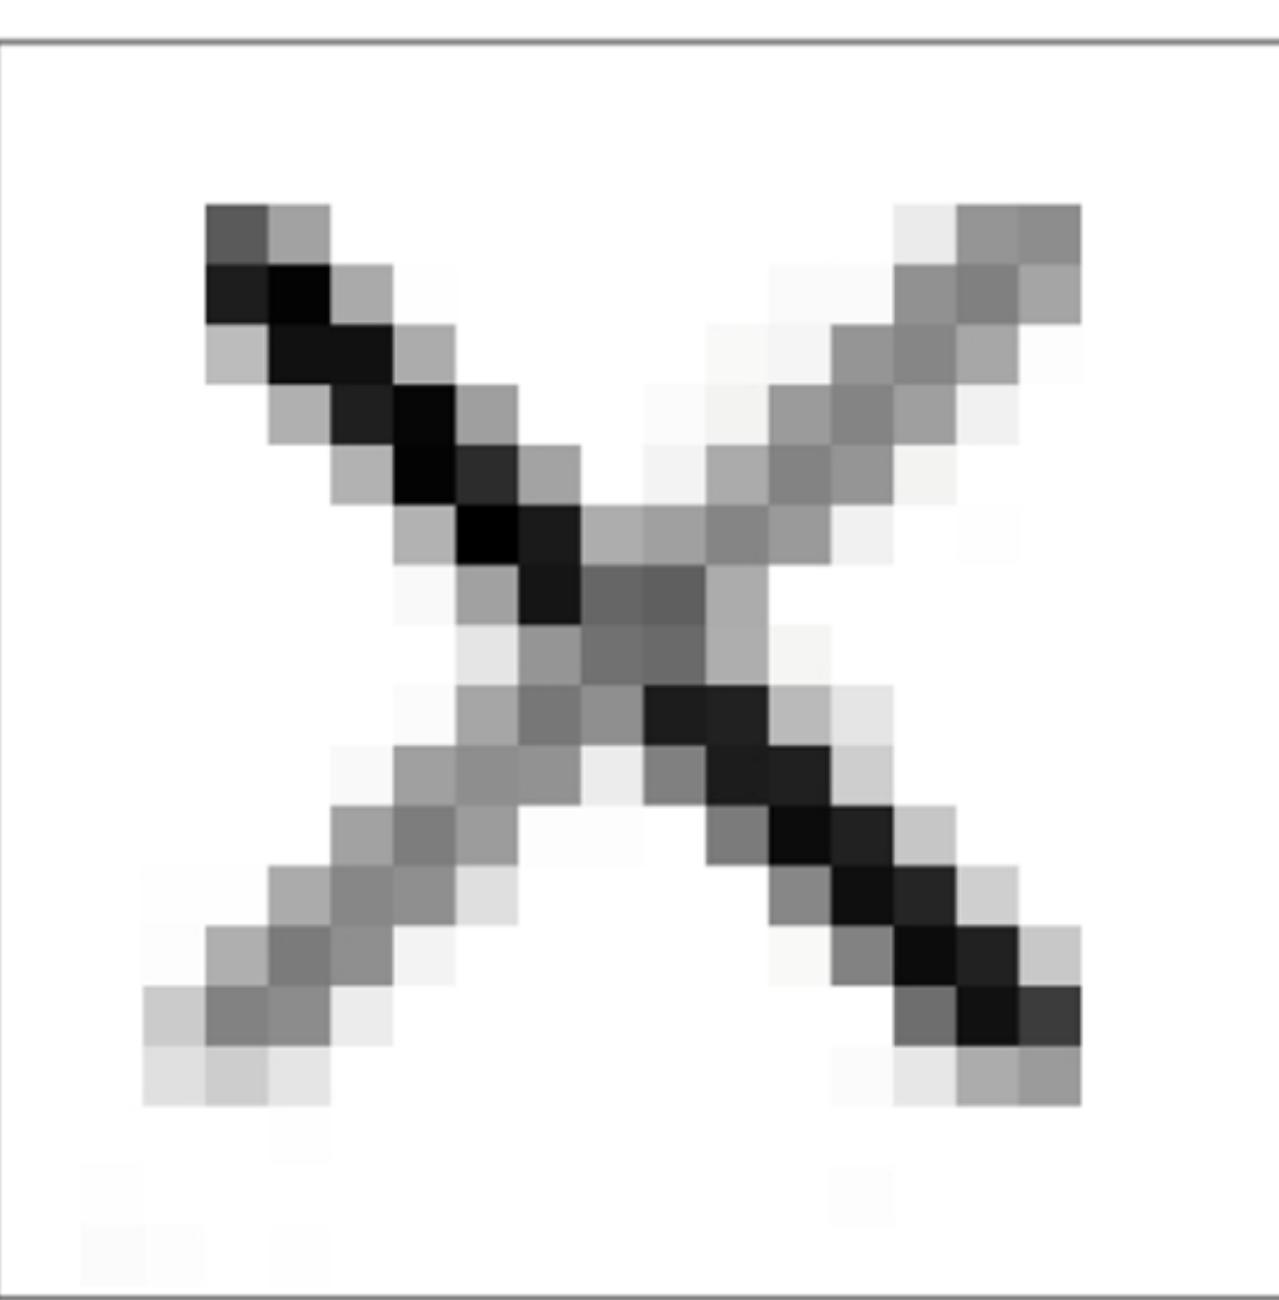


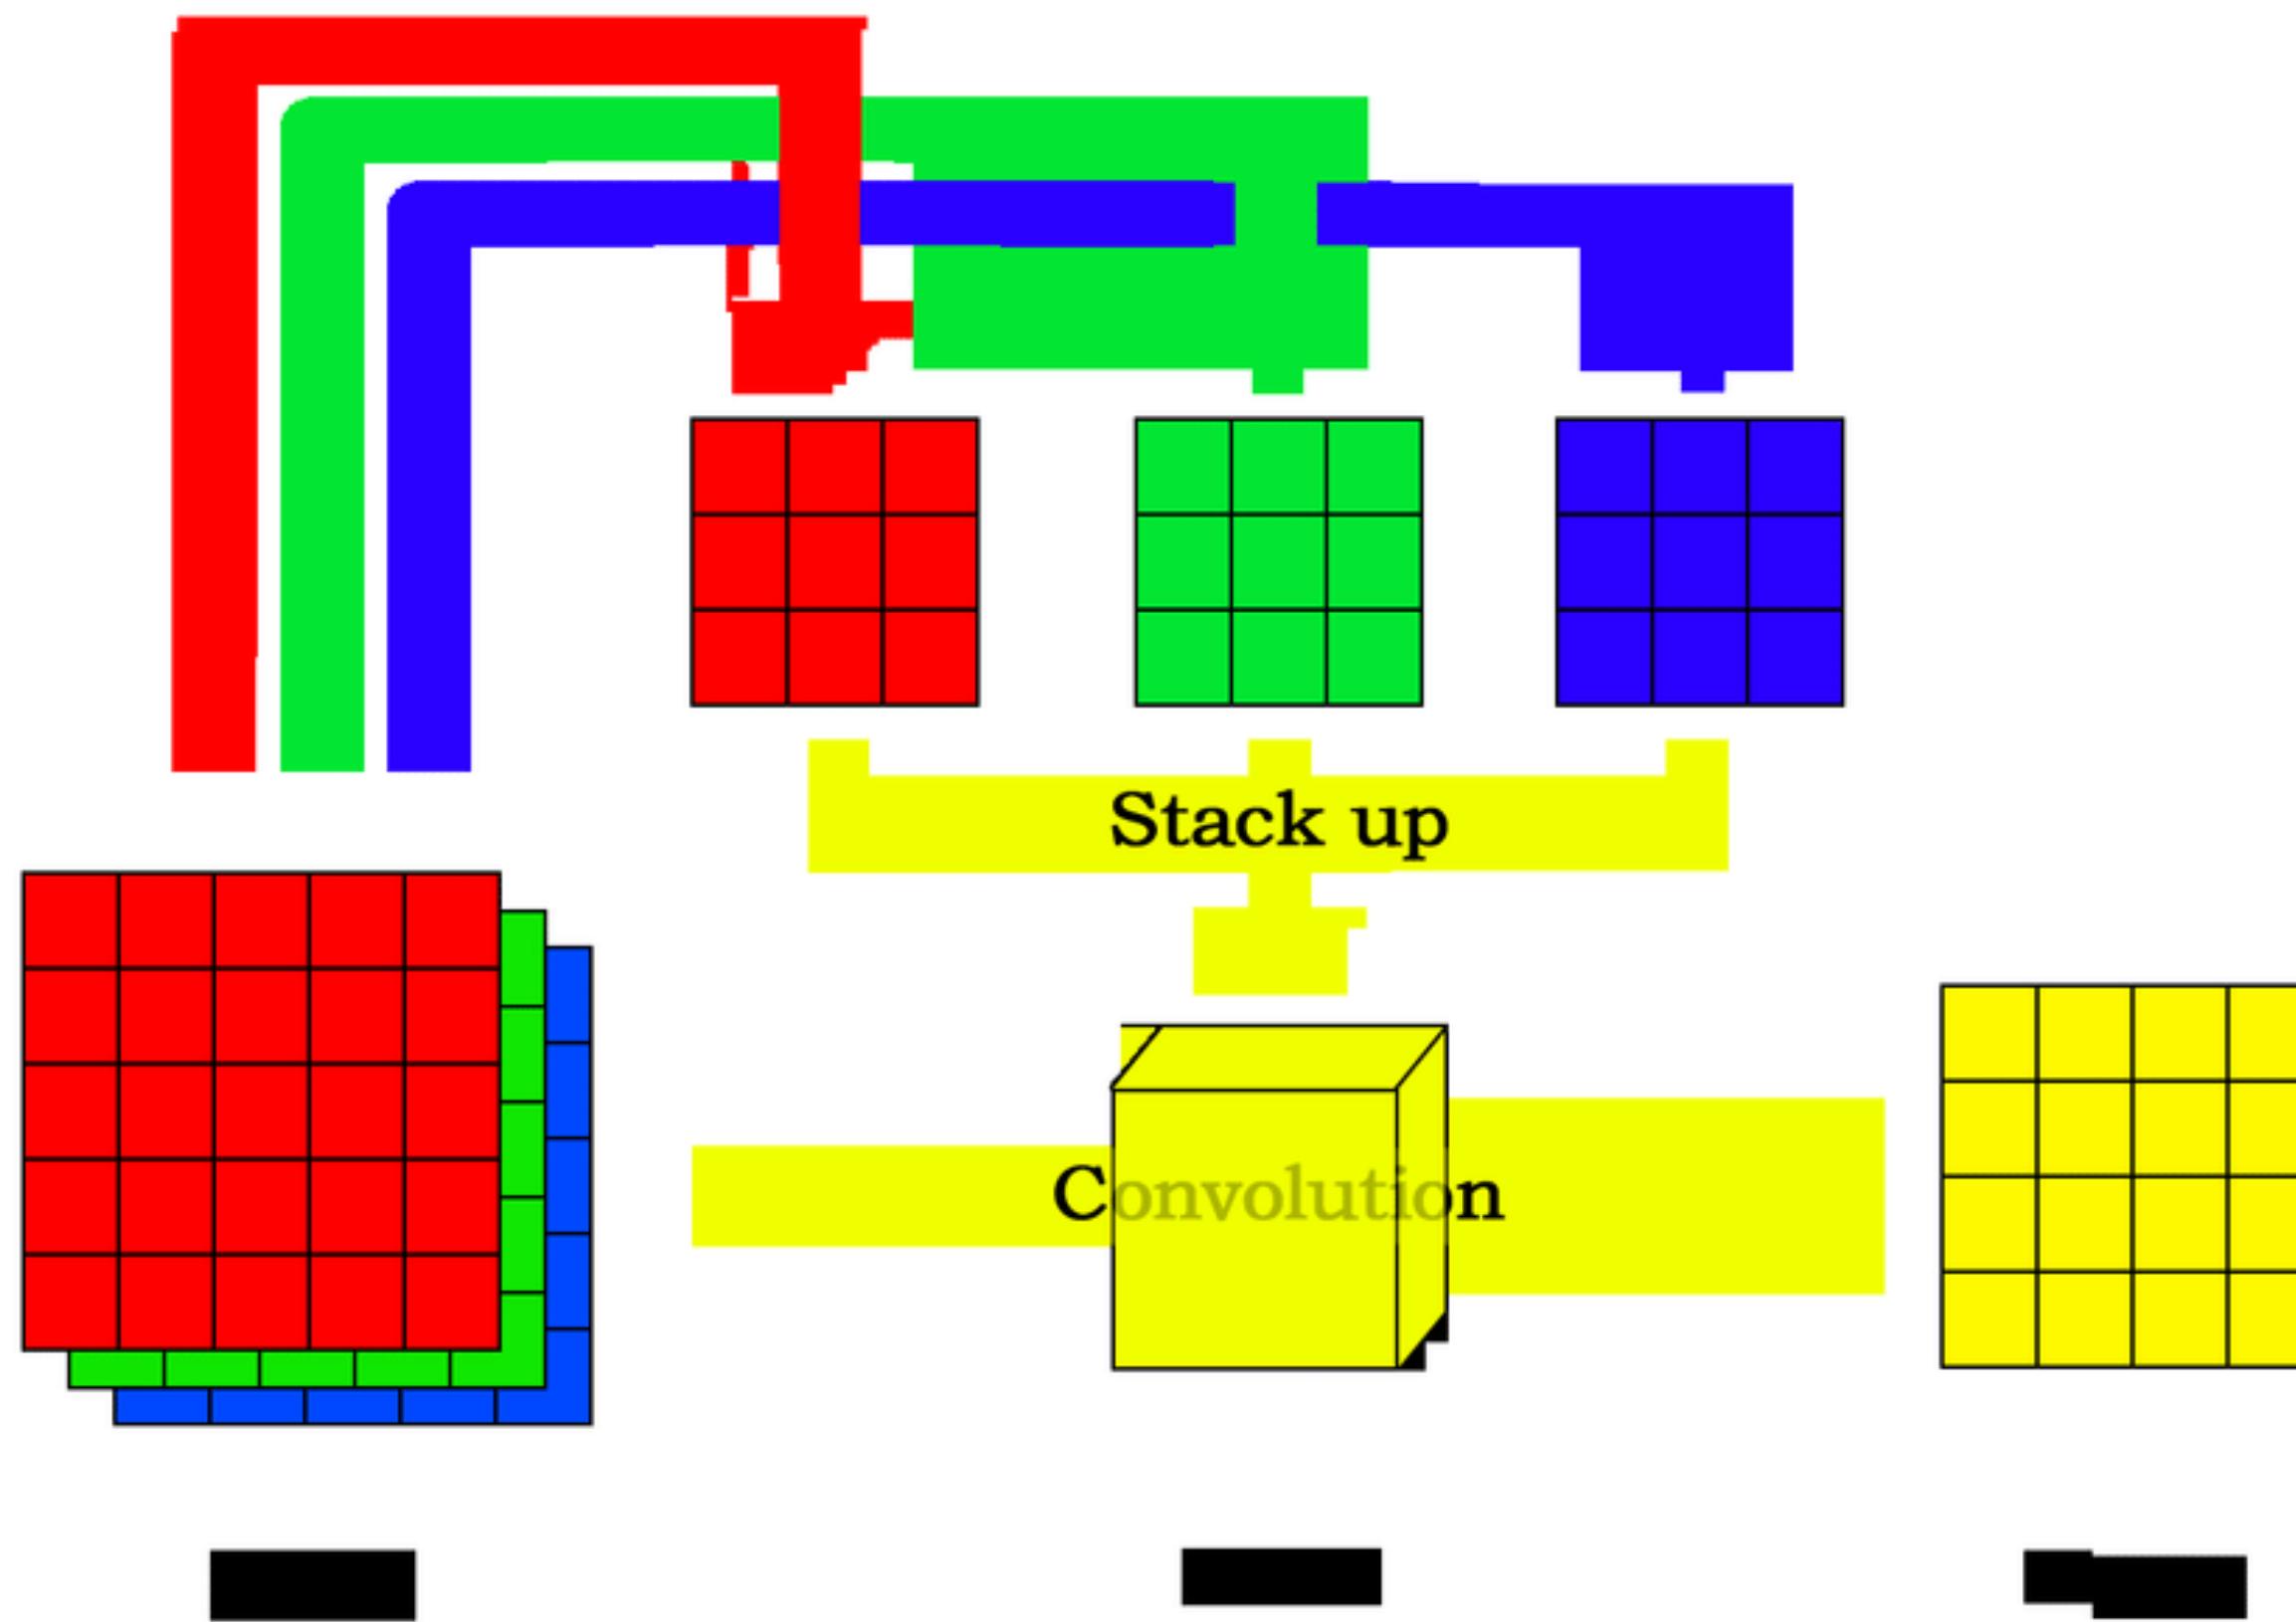
**Convolutional Neural Networks (CNNs) are a category of Neural Network that have proven very effective in areas such as image recognition and classification.**





# Input image

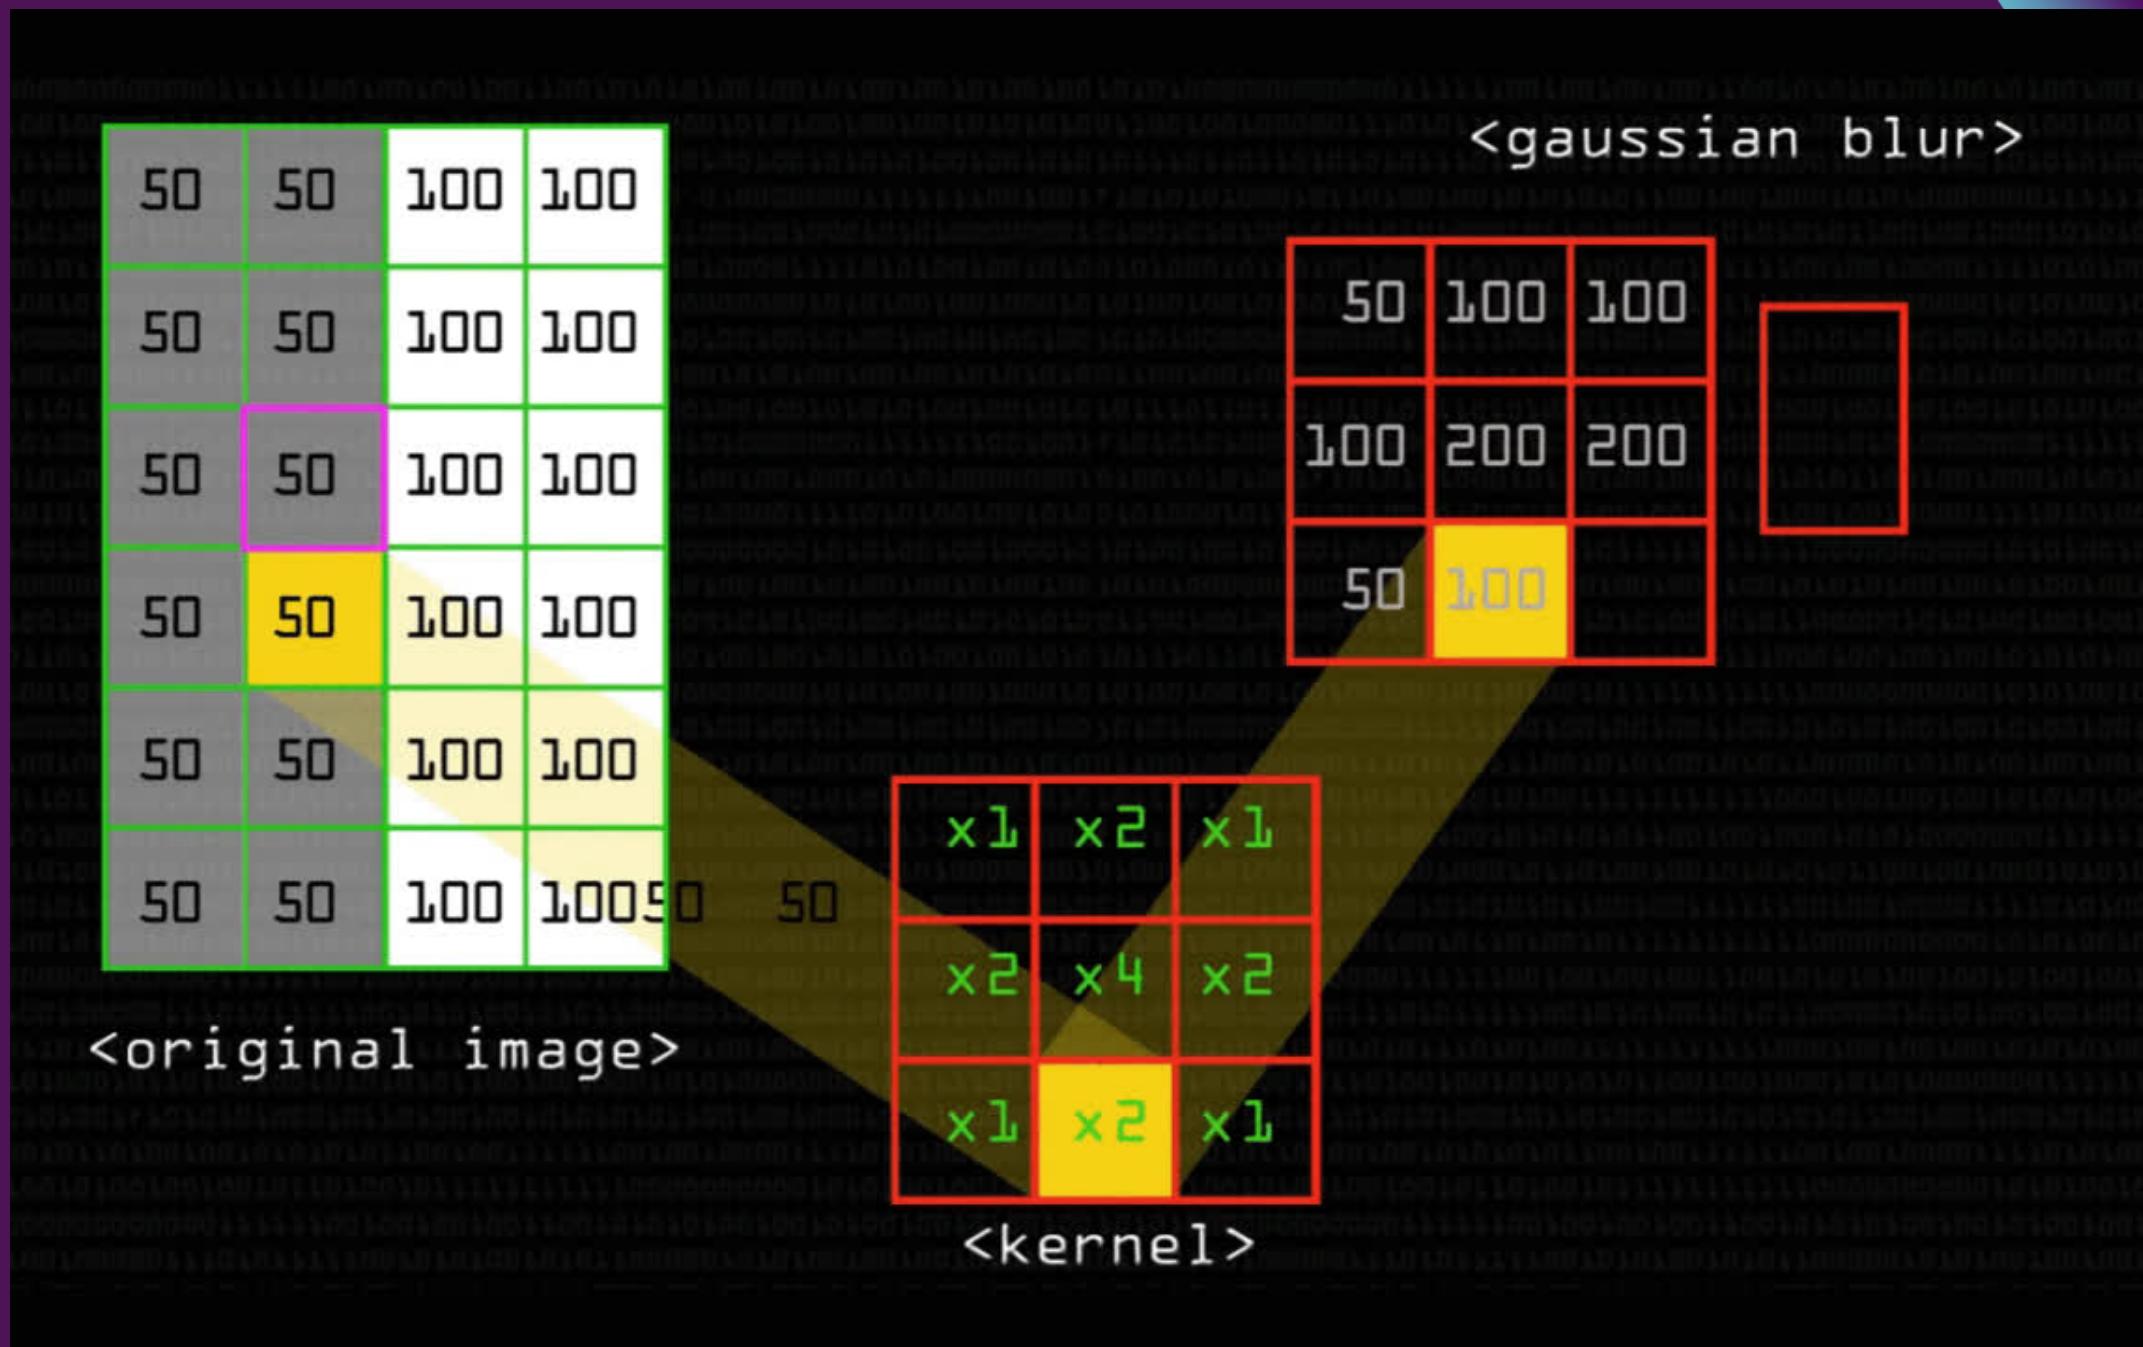




# Convolutional layers

Convolutional layers are a type of neural network layer primarily used in deep learning models for processing data that has a grid-like topology, such as images.

These layers perform a mathematical operation called convolution



# Convolutional layers

- **Convolutional Layer:** The primary purpose of Convolution is to extract features from the input image. Convolution preserves the spatial relationship between input by learning input features.

$$\text{Convolution layer output } y = \sum_i^n x * f$$

1	1	1	0	0
0	1	1	1	0
0	0	-1	1	1
0	0	-1	1	0
0	1	-1	0	0

\*

1	0	1
0	1	0
1	0	1

=

2	3	2
0	2	1
-2	1	0

Input Image(x)  
(nxn)

Filter  
(fxf)

Output Image (y)  
(n-f+1) x (n-f+1)



# Convolutional layers

- Kernel is nothing but a filter used to extract the features from the images
  - The kernel is a matrix that moves over the input data, performs the dot product with the sub-region of input data, and sets the output as the matrix of dot products.
  - Kernel moves on the input data by the stride value
- Strides: The number of pixels shifts over the input matrix by the filter. When the stride is 1 then we move the filters to 1 pixel at a time.



Reference: <https://www.quora.com/How-are-convolutional-filters-kernels-initialized-and-learned-in-a-convolutional-neural-network-CNN>

# Pooling

**Pooling layers would reduce the number of parameters when the inputs are too large.**

**Pooling also called down sampling which reduces the dimensionality of each map but retains important information.**

**There are three types of pooling namely, Max Pooling, Average Pooling, Sum Pooling.**



**Max Pooling**

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

**2 x 2 pool size**

100	184
12	45

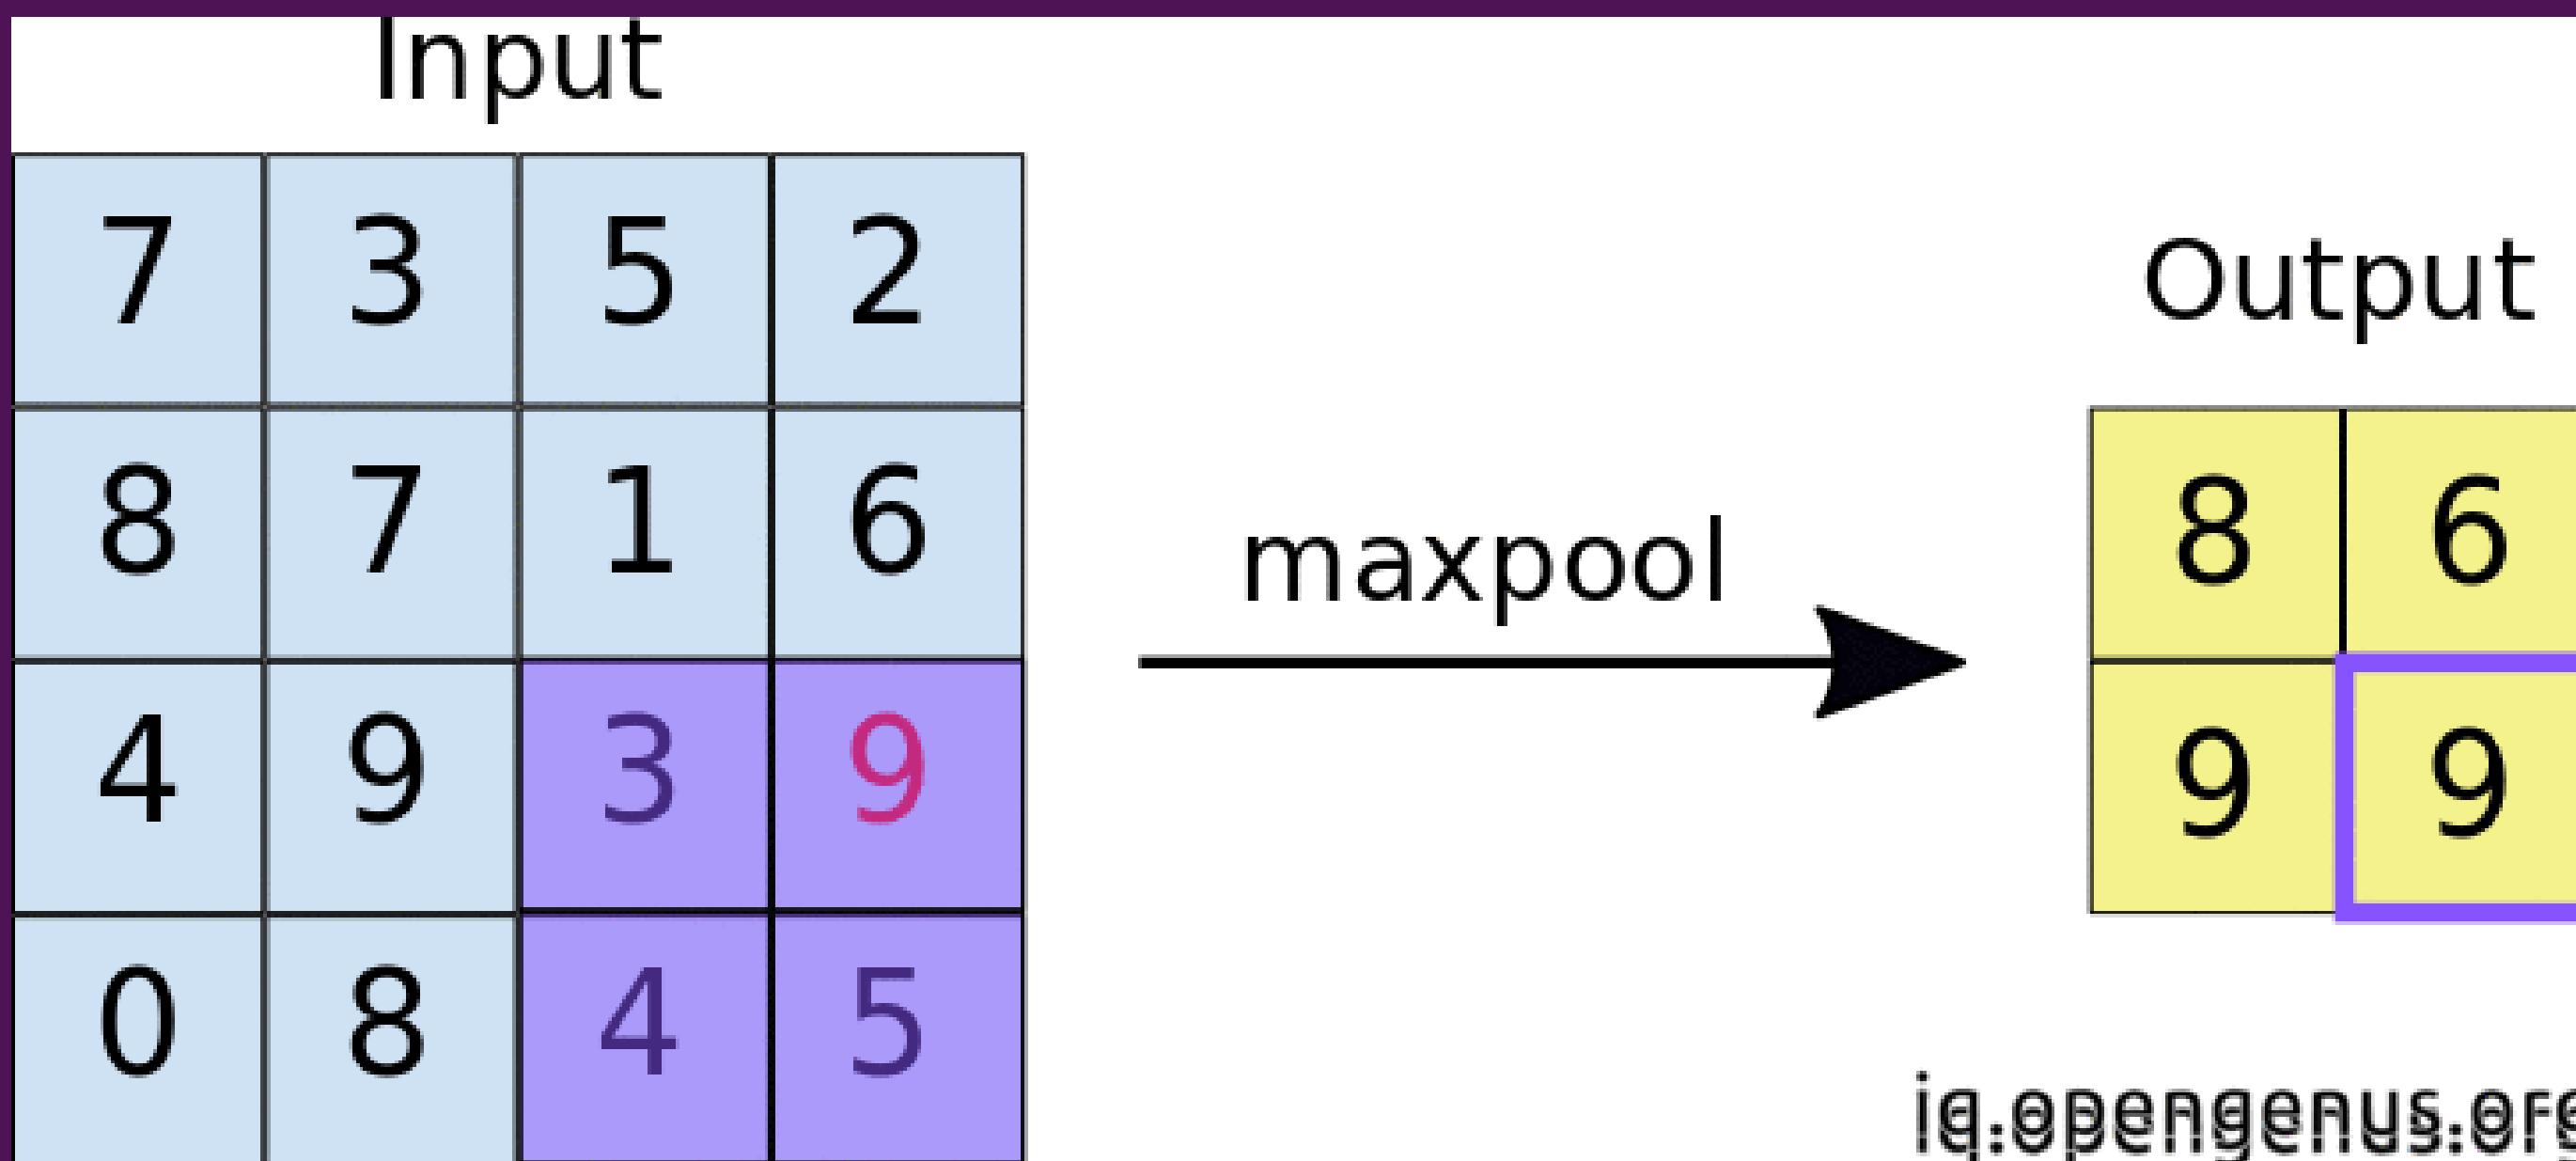
**Average Pooling**

31	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

**2 x 2 pool size**

36	80
12	15

# Pooling



# Flatten Layer

The Flatten layer is introduced to convert the multi-dimensional tensor into a one-dimensional array. It essentially unravels or flattens the spatial structure of the data while retaining the learned features.

## Flatten Layer in Keras

1	1	0
4	2	1
0	2	1

Pooled  
Feature Map

Flattening

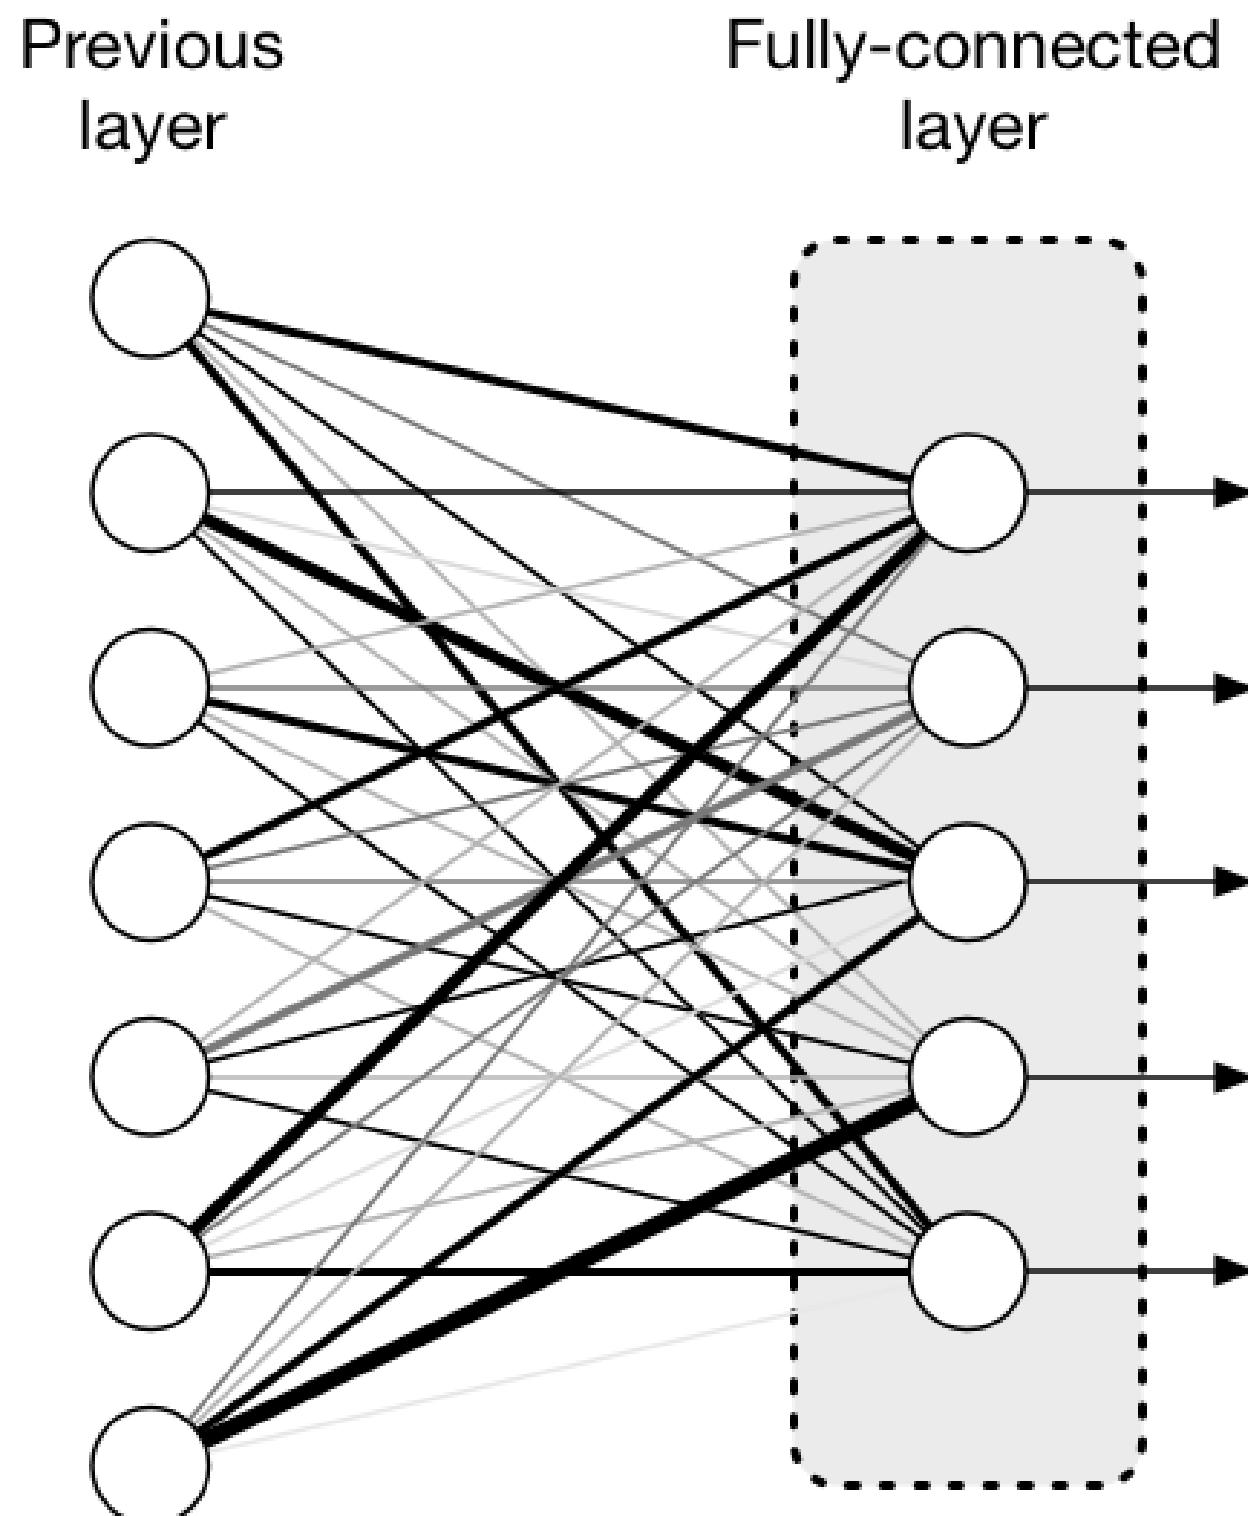
1
1
0
4
2
1
0
2
1



# Fully connected layer

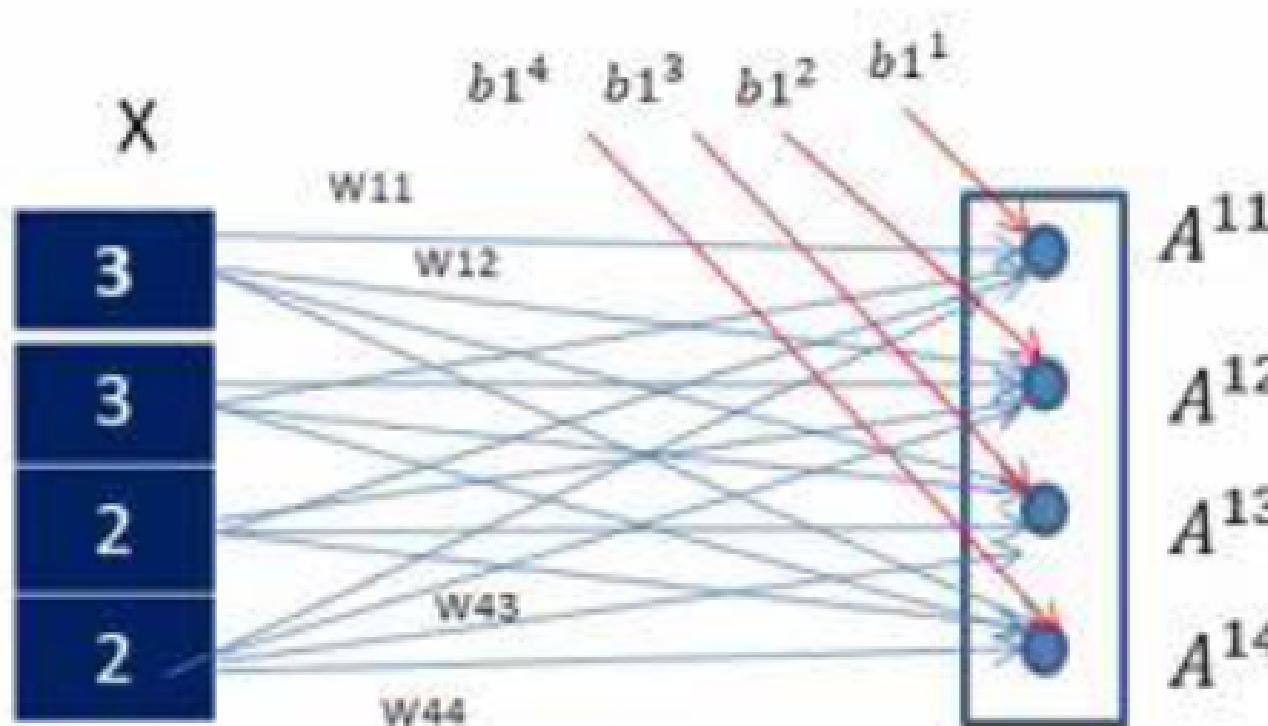
**Dense (Fully Connected) Layers** are typical in the later stages of a CNN. These layers can learn complex relationships across different parts of the image and produce the final output of the network

These layers are responsible for making predictions based on the high-level features learned by the convolutional layers



# Fully connected layer

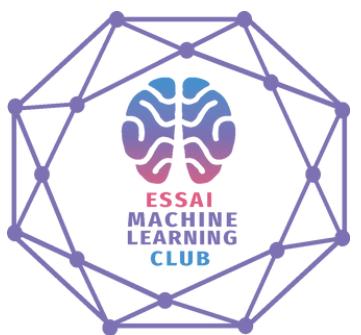
## Fully Connected layer with ReLU



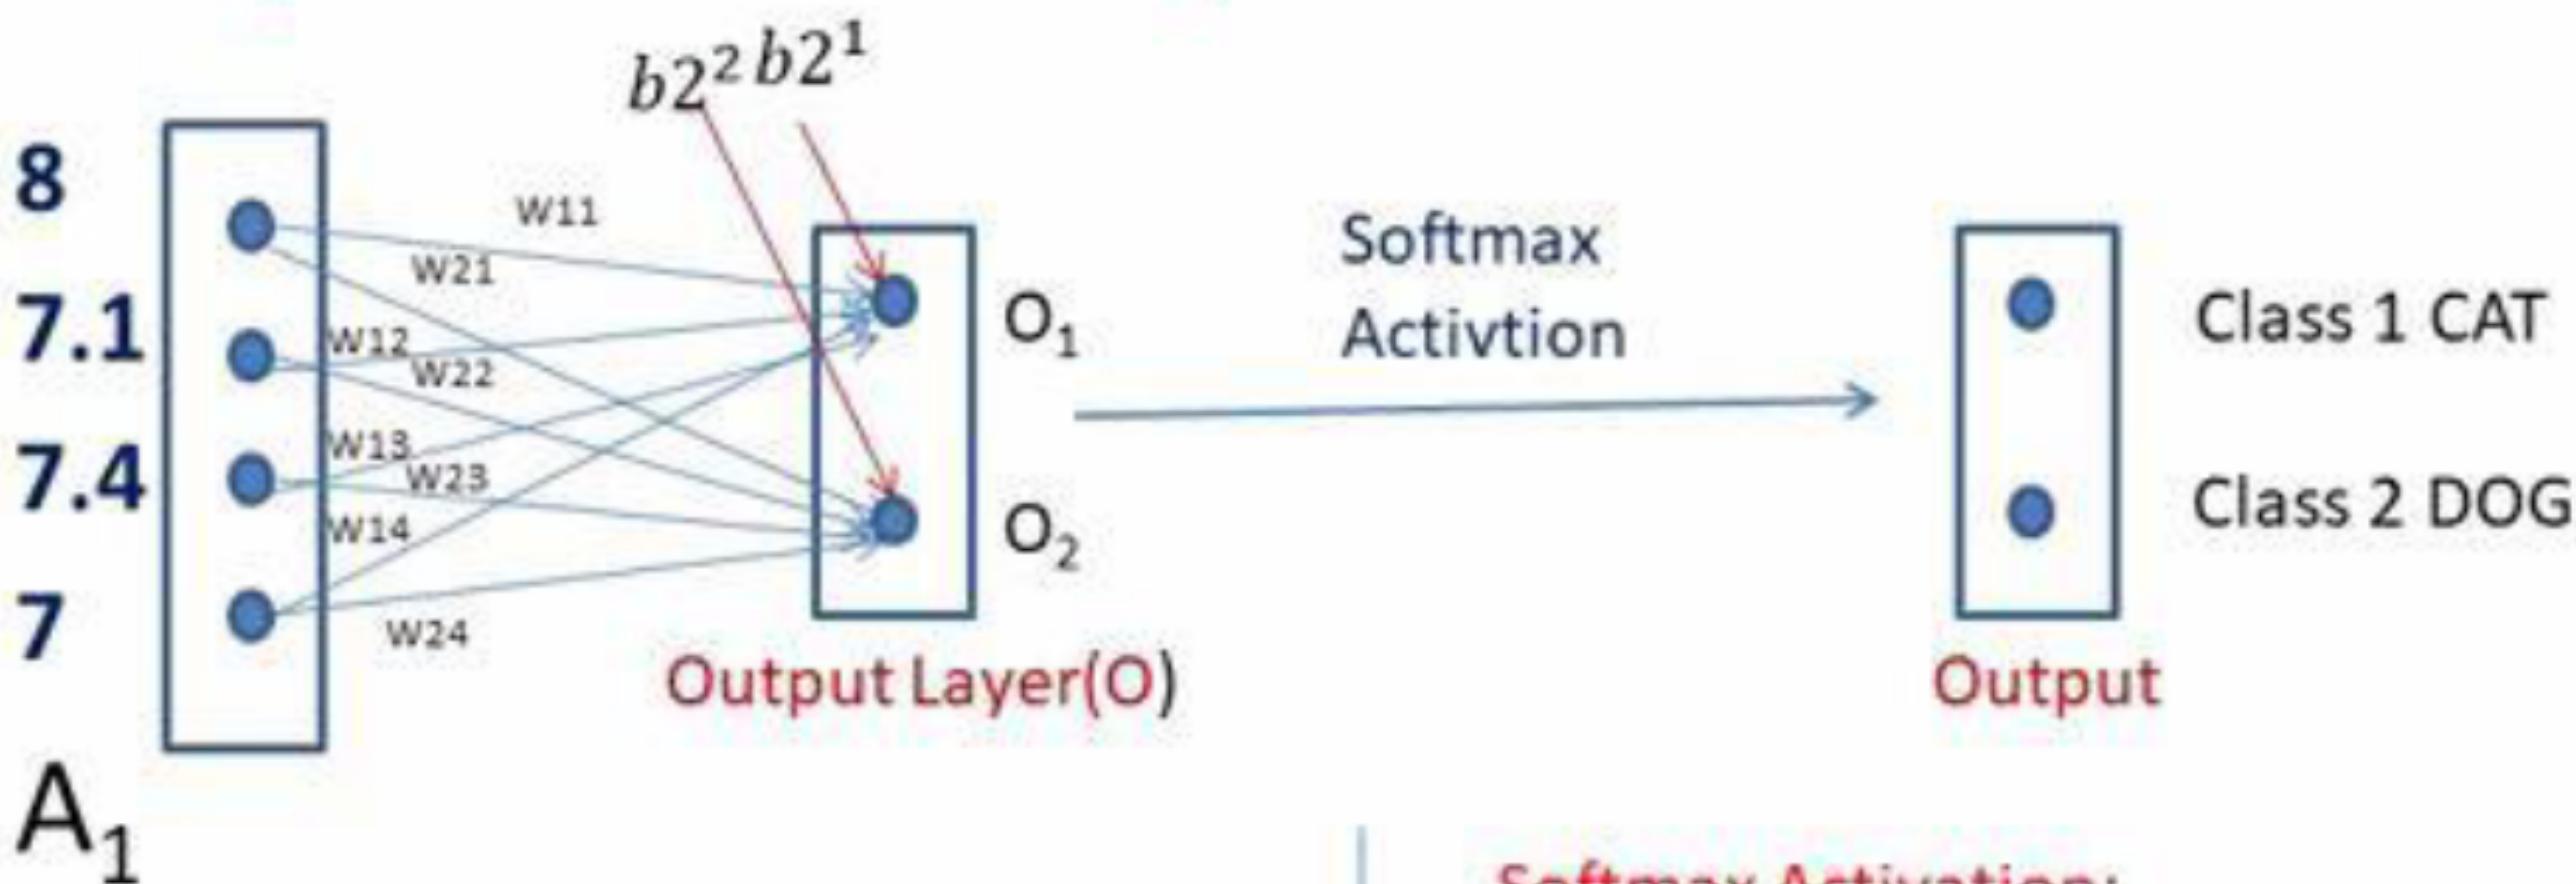
$$A_1 = W_1 X + b_1$$

$$\begin{bmatrix} A_{11} \\ A_{12} \\ A_{13} \\ A_{14} \end{bmatrix} = \begin{bmatrix} 1.0 & 1.0 & 0.2 & 0.8 \\ 1.0 & 0.5 & 0.5 & 0.8 \\ 0.8 & 1.0 & 0.2 & 0.8 \\ 0.5 & 0.5 & 1.0 & 1.0 \end{bmatrix} \begin{bmatrix} 3 \\ 3 \\ 2 \\ 2 \end{bmatrix} + \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \\ 0.1 \end{bmatrix} = \begin{bmatrix} 8 \\ 7.1 \\ 7.4 \\ 7 \end{bmatrix}$$

$$\text{ReLU} \begin{bmatrix} 8 \\ 7.1 \\ 7.4 \\ 7 \end{bmatrix} = \begin{bmatrix} 8 \\ 7.1 \\ 7.4 \\ 7 \end{bmatrix}$$



# Fully Connected layer & Softmax Activation



$$O = W_2 A_1 + b_2$$

$$\begin{bmatrix} O_1 \\ O_2 \end{bmatrix} = \begin{bmatrix} 0.2 & 0.1 & 0.5 & 0.5 \\ 0.1 & 0.5 & 0.1 & 0.2 \end{bmatrix} \begin{bmatrix} 8 \\ 7.1 \\ 7.4 \\ 7 \end{bmatrix} + \begin{bmatrix} 0.1 \\ 0.1 \end{bmatrix}$$
$$= \begin{bmatrix} 9.61 \\ 6.49 \end{bmatrix}$$

Softmax Activation:

$$\sigma(O_i) = \frac{e^{o_i}}{\sum e^{o_i}}$$

$$\sigma(O_1) = \frac{e^{o_1}}{e^{o_1} + e^{o_2}} = \frac{14913.17}{14913.17 + 658.52} = 0.95$$

$$\sigma(O_2) = \frac{e^{o_2}}{e^{o_1} + e^{o_2}} = \frac{658.52}{14913.17 + 658.52} = 0.04$$



**Transfer Learning** : Transfer learning involves taking a pre-trained model (usually trained on a large dataset) and using it as a starting point for a new task or dataset.

examples

VGG (Visual Geometry Group)   ResNet (Residual Network)   MobileNet   DenseNet   EfficientNet

**Fine-tuning** : Fine-tuning is a specific form of transfer learning where instead of keeping all the pre-trained model parameters frozen, you unfreeze some or all of the layers and allow them to be further trained on the new dataset.



# CS231n Convolutional Neural Networks for Visual Recognition

Course materials and notes for Stanford class CS231n: Convolutional Neural Networks for Visual Recognition.

[github](#) [in](#)



## A Comprehensive Survey on Transfer Learning

Transfer learning aims at improving the performance of target learners on target domains by transferring the knowledge contained in different but related source domains. In this way, the...



## What is the difference between Transfer Learning vs Fine Tuning...

I finished my undergraduate thesis in transfer learning, keeping the most of the CNN layer freeze...

R<sup>6</sup> ResearchGate



TensorFlow



Keras

PyTorch



# Deep Learning Aproach For Safty Roads

01 Introduction to  
Computer Vision



[Yahia Chammami](#)